

# X86-64

Cauchy(pqy7172@gmail.com)

December 28, 2025

## Contents

<b>1</b>		<b>2</b>
1.1	.....	2
1.2	.....	3
<b>2</b>		<b>26</b>
X86-64		
=> e1000_intr		
=> __handle_irq_event_percpu		
=> handle_irq_event		
=> handle_fastehoi_irq		
=> __common_interrupt		
=> common_interrupt		
=> asm_common_interrupt		
=> e1000_clean_rx_irq		
=> e1000_clean		
=> __napi_poll		
=> net_rx_action		
=> handle_softirqs		
=> __irq_exit_rcu		
=> common_interrupt		
=> asm_common_interrupt		
=> finish_task_switch.isra.0		
=> __schedule		
=> schedule		
=> worker_thread		
=> kthread		
=> ret_from_fork		
=> ret_from_fork_asm		

# 1

X86-64	idt_setup_apic_and_irq_gates	irq_entries_start	X86-64
--------	------------------------------	-------------------	--------

## 1.1

```
.align IDT_ALIGN
SYM_CODE_START(irq_entries_start)
    vector=FIRST_EXTERNAL_VECTOR
    .rept NR_EXTERNAL_VECTORS
    UNWIND_HINT_IRET_REGS
0 :
    ENDBR
    .byte 0x6a, vector
    jmp asm_common_interrupt
/* Ensure that the above is IDT_ALIGN bytes max */
    .fill 0b + IDT_ALIGN - ., 1, 0xcc
    vector = vector+1
    .endr
SYM_CODE_END(irq_entries_start)

NR_EXTERNAL_VECTORS FIRST_EXTERNAL_VECTOR

/*
 * Posted interrupt notification vector for all device MSIs delivered to
 * the host kernel.
*/
#define POSTED_MSI_NOTIFICATION_VECTOR 0xeb

/*
 * IDT vectors usable for external interrupt sources start at 0x20.
 * (0x80 is the syscall vector, 0x30-0x3f are for ISA)
*/
#define FIRST_EXTERNAL_VECTOR 0x20

#ifndef CONFIG_X86_LOCAL_APIC
#define FIRST_SYSTEM_VECTOR POSTED_MSI_NOTIFICATION_VECTOR
#else
#define FIRST_SYSTEM_VECTOR NR_VECTORS
#endif
#define NR_EXTERNAL_VECTORS (FIRST_SYSTEM_VECTOR - FIRST_EXTERNAL_VECTOR)
```

X86-64	0x20	PIC APIC Local APIC	CPU
--------	------	---------------------	-----

```

x86-64      CONFIG_X86_LOCAL_APIC FIRST_SYSTEM_VECTOR POSTED_MSI_NOTIFICATION
0xEB          MSI Message Signaled Interrupts           vCPU   vCPU   APICv APIC PI
Exit          I/O                                POSTED_MSI_NOTIFICATION_VECTOR PI    vCP
Exit vCPU      VM-Exit
[0x20, 0xeb] 203      0x20     CPU    0     Page Fault NMI      0x20
.byte push    push      push      127     push+   5 1 push  4      32     IDT
entry stub    IDT_ALIGN
vmlinux      irq_entries_start

fffffffff81e00230 <irq_entries_start>:
fffffffff81e00230:  f3 0f 1e fa      endbr64
fffffffff81e00234:  6a 20      push   $0x20
fffffffff81e00236:  e9 05 13 00 00  jmp    ffffffff81e01540 <asm_common_interrup
fffffffff81e0023b:  cc      int3
fffffffff81e0023c:  cc      int3
fffffffff81e0023d:  cc      int3
fffffffff81e0023e:  cc      int3
fffffffff81e0023f:  cc      int3
fffffffff81e00240:  f3 0f 1e fa      endbr64
fffffffff81e00244:  6a 21      push   $0x21
fffffffff81e00246:  e9 f5 12 00 00  jmp    ffffffff81e01540 <asm_common_interrup
fffffffff81e0024b:  cc      int3
fffffffff81e0024c:  cc      int3
fffffffff81e0024d:  cc      int3
fffffffff81e0024e:  cc      int3
fffffffff81e0024f:  cc      int3

push          endbr64      X86_KERNEL_IBT
Control Protection Exception
.align IDT_ALIGN           Intel      IBT(Intel CET Control-flow En-
forment Technology)           16
.rept          NR_EXTERNAL_VECTORS .


```

## 1.2

```

        idt_setup_apic_and_irq_gates

/**
 * idt_setup_apic_and_irq_gates - Setup APIC/SMP and normal interrupt gates
 */
void __init idt_setup_apic_and_irq_gates(void)
{
    int i = FIRST_EXTERNAL_VECTOR;
    void *entry;

    idt_setup_from_table(idt_table, apic_idts, ARRAY_SIZE(apic_idts), true);

```

```

for_each_clear_bit_from(i, system_vectors, FIRST_SYSTEM_VECTOR) {
    entry = irq_entries_start + IDT_ALIGN * (i - FIRST_EXTERNAL_VECTOR);
    set_intr_gate(i, entry);
}

#ifndef CONFIG_X86_LOCAL_APIC
for_each_clear_bit_from(i, system_vectors, NR_VECTORS) {
/*
 * Don't set the non assigned system vectors in the
 * system_vectors bitmap. Otherwise they show up in
 * /proc/interrupts.
 */
entry = spurious_entries_start + IDT_ALIGN * (i - FIRST_SYSTEM_VECTOR);
set_intr_gate(i, entry);
}
#endif
/* Map IDT into CPU entry area and reload it. */
idt_map_in_cea();
load_idt(&idt_descr);

/* Make the IDT table read only */
set_memory_ro((unsigned long)&idt_table, 1);

idt_setup_done = true;
}

apic_idts      idt_table          set_intr_gate idt_table      idt_map_i
idt_setup_from_table

arch/x86/kernel/idt.c
static __init void
idt_setup_from_table(gate_desc *idt, const struct idt_data *t, int size, bool sys)
{
    gate_desc desc;

    for (; size > 0; t++, size--) {
        idt_init_desc(&desc, t);
        write_idt_entry(idt, t->vector, &desc);
        if (sys)
            set_bit(t->vector, system_vectors);
    }
}

/* apic_idts
arch/x86/kernel/idt.c
*/

```

```

* The APIC and SMP idt entries
*/
static const __initconst struct idt_data apic_idts[] = {
#endif CONFIG_SMP
    INTG(RESCHEDULE_VECTOR,    asm_sysvec_reschedule_ipi),
    INTG(CALL_FUNCTION_VECTOR, asm_sysvec_call_function),
    INTG(CALL_FUNCTION_SINGLE_VECTOR, asm_sysvec_call_function_single),
    INTG(REBOOT_VECTOR,       asm_sysvec_reboot),
#endif

#ifndef CONFIG_X86_THERMAL_VECTOR
    INTG(THERMAL_APIC_VECTOR,  asm_sysvec_thermal),
#endif

#ifndef CONFIG_X86_MCE_THRESHOLD
    INTG(THRESHOLD_APIC_VECTOR,  asm_sysvec_threshold),
#endif

#ifndef CONFIG_X86_MCE_AMD
    INTG(DEFERRED_ERROR_VECTOR,  asm_sysvec_deferred_error),
#endif

#ifndef CONFIG_X86_LOCAL_APIC
    INTG(LOCAL_TIMER_VECTOR,    asm_sysvec_apic_timer_interrupt),
    INTG(X86_PLATFORM_IPI_VECTOR,  asm_sysvec_x86_platform_ipi),
#endif IS_ENABLED(CONFIG_KVM)
    INTG(POSTED_INTR_VECTOR,    asm_sysvec_kvm_posted_intr_ipi),
    INTG(POSTED_INTR_WAKEUP_VECTOR,  asm_sysvec_kvm_posted_intr_wakeup_ipi),
    INTG(POSTED_INTR_NESTED_VECTOR,  asm_sysvec_kvm_posted_intr_nested_ipi),
#endif
#ifndef CONFIG_IRQ_WORK
    INTG(IRQ_WORK_VECTOR,      asm_sysvec_irq_work),
#endif
    INTG(SPURIOUS_APIC_VECTOR,  asm_sysvec_spurious_apic_interrupt),
    INTG(ERROR_APIC_VECTOR,    asm_sysvec_error_interrupt),
#endif CONFIG_X86_POSTED_MSI
    INTG(POSTED_MSI_NOTIFICATION_VECTOR,  asm_sysvec_posted_msi_notification),
#endif
#endif
};

asm_sysvec_reschedule_ipi
DECLARE_IDTENTRY(RESCHEDULE_VECTOR,    sysvec_reschedule_ipi);
DECLARE_IDTENTRY      .c  .S      .c
arch/x86/include/asm/idtentry.h

```

```

/***
 * DECLARE_IDTENTRY - Declare functions for simple IDT entry points
 *           No error code pushed by hardware
 * @vector: Vector number (ignored for C)
 * @func: Function name of the entry point
 *
 * Declares four functions:
 * - The ASM entry point: asm_##func
 * - The XEN PV trap entry point: xen_##func (maybe unused)
 * - The C handler called from the FRED event dispatcher (maybe unused)
 * - The C handler called from the ASM entry point
 *
 * Note: This is the C variant of DECLARE_IDTENTRY(). As the name says it
 * declares the entry points for usage in C code. There is an ASM variant
 * as well which is used to emit the entry stubs in entry_32/64.S.
 */
#define DECLARE_IDTENTRY(vector, func)      \
    asmlinkage void asm_##func(void);      \
    asmlinkage void xen_asm_##func(void);   \
    void fred_##func(struct pt_regs *regs); \
    __visible void func(struct pt_regs *regs)

        asm_sysvec_reschedule_ipi          .S  asm_sysvec_reschedule_ipi

arch/x86/include/asm/idtentry.h

#else /* !_ASSEMBLER_ */

/*
 * The ASM variants for DECLARE_IDTENTRY() which emit the ASM entry stubs.
 */
#define DECLARE_IDTENTRY(vector, func)      \
    idtentry vector asm_##func func has_error_code=0

idtentry arch/x86/entry/entry_64.S      asm_##func  asm_sysvec_reschedule_ipi  idtentry  ent

fffffffff81e01630 <asm_sysvec_reschedule_ipi>:
fffffffff81e01630:  f3 0f 1e fa          endbr64
fffffffff81e01634:  90                   nop
fffffffff81e01635:  90                   nop
fffffffff81e01636:  90                   nop
fffffffff81e01637:  fc                   cld
fffffffff81e01638:  6a ff               push   $0xfffffffffffffff
fffffffff81e0163a:  e8 f1 05 00 00       call   ffffffff81e01c30 <error_entry>
fffffffff81e0163f:  48 89 c4             mov    %rax,%rsp
fffffffff81e01642:  48 89 e7             mov    %rsp,%rdi
fffffffff81e01645:  e8 36 35 ef ff       call   ffffffff81cf4b80 <sysvec_reschedule_i

```

```

ffffffffff81e0164a:      e9 21 07 00 00      jmp      ffffffff81e01d70 <error_return>
ffffffffff81e0164f:      90                      nop

asm_sysvec_reboot

ffffffffff81e01650 <asm_sysvec_reboot>:
ffffffffff81e01650:      f3 0f 1e fa          endbr64
ffffffffff81e01654:      90                      nop
ffffffffff81e01655:      90                      nop
ffffffffff81e01656:      90                      nop
ffffffffff81e01657:      fc                      cld
ffffffffff81e01658:      6a ff          push    $0xfffffffffffffff
ffffffffff81e0165a:      e8 d1 05 00 00      call    ffffffff81e01c30 <error_entry>
ffffffffff81e0165f:      48 89 c4          mov     %rax,%rsp
ffffffffff81e01662:      48 89 e7          mov     %rsp,%rdi
ffffffffff81e01665:      e8 86 34 ef ff      call    ffffffff81cf4af0 <sysvec_reboot>
ffffffffff81e0166a:      e9 01 07 00 00      jmp    ffffffff81e01d70 <error_return>
ffffffffff81e0166f:      90                      nop

error_entry      sysvec_reschedule_ipi/sysvec_reboot      error_return      call
sysvec_reschedule_ipi

idtentry->idtentry_body->call \cfunc

cfunc idtentry  cfunc      DECLARE_IDTENTRY sysvec_reschedule_ipi
DECLARE_IDTENTRY(RESCHEDULE_VECTOR,   sysvec_reschedule_ipi);

call sysvec_reschedule_ipi

arch/x86/kernel/smp.c

DEFINE_IDTENTRY_SYSVEC_SIMPLE(sysvec_reschedule_ipi)
{
    apic_eoi();
    trace_reschedule_entry(RESCHEDULE_VECTOR);
    inc_irq_stat(irq_resched_count);
    scheduler_ipi();
    trace_reschedule_exit(RESCHEDULE_VECTOR);
}

arch/x86/include/asm/idtentry.h

#define DEFINE_IDTENTRY_SYSVEC_SIMPLE(func) \
static __always_inline void __##func(struct pt_regs *regs); \
\ \
static __always_inline void instr_##func(struct pt_regs *regs) \
{ \
    \
    __irq_enter_raw(); \
}

```

```

__##func (regs);      \
__irq_exit_raw();    \
}                   \
\                  \
__visible noinstr void func(struct pt_regs *regs)  \
{                 \
    irqentry_state_t state = irqentry_enter(regs);  \
    \
    kvm_set_cpu_l1tf_flush_l1d();                      \
    instrumentation_begin();                           \
    instr_##func (regs);                            \
    instrumentation_end();                          \
    irqentry_exit(regs, state);                    \
}                   \
\                  \
void fred_##func(struct pt_regs *regs)      \
{                 \
    instr_##func (regs);                            \
}                   \
\                  \
static __always_inline void __##func(struct pt_regs *regs)
    __##func

{
    apic_eoi();
    trace_reschedule_entry(RESCHEDULE_VECTOR);
    inc_irq_stat(irq_resched_count);
    scheduler_ipi();
    trace_reschedule_exit(RESCHEDULE_VECTOR);
}

identity           /          cfunc   sysvec_reschedule_ipi

DECLARE_IDTENTRY_SYSVEC(REBOOT_VECTOR,   sysvec_reboot);
/           XXX       funcxxx
DECLARE_IDTENTRY_SYSVEC(XXX,  funcxxx)

DEFINE_IDTENTRY_SYSVEC_SIMPLE
DEFINE_IDTENTRY_SYSVEC_SIMPLE(funcxxx)
sysvec_reschedule_ipi     scheduler_ipi
DECLARE_IDTENTRY_SYSVEC->DECLARE_IDTENTRY->idtentry

asm_sysvec_reschedule_ipi/asm_sysvec_reboot      addr  INTG

```

```

#define G(_vector, _addr, _ist, _type, _dpl, _segment) \
{ \
    .vector  = _vector, \
    .bits.ist = _ist, \
    .bits.type = _type, \
    .bits.dpl = _dpl, \
    .bits.p  = 1, \
    .addr   = _addr, \
    .segment = _segment, \
}

/* Interrupt gate */
#define INTG(_vector, _addr) \
G(_vector, _addr, DEFAULT_STACK, GATE_INTERRUPT, DPL0, __KERNEL_CS)

        idt_data::addr apic_idts           apic_idts     0-31 cpu    32-NR_EXTERNAL_VECTORS
        idt_setup_apic_and_irq_gates->idt_setup_from_table

static __init void
idt_setup_from_table(gate_desc *idt, const struct idt_data *t, int size, bool sys)
{
    gate_desc desc;

    for (; size > 0; t++, size--) {
        idt_init_desc(&desc, t);
        write_idt_entry(idt, t->vector, &desc);
        if (sys)
            set_bit(t->vector, system_vectors);
    }
}

apic_idts           idt_init_desc

static inline void idt_init_desc(gate_desc *gate, const struct idt_data *d)
{
    unsigned long addr = (unsigned long) d->addr;

    gate->offset_low = (u16) addr;
    gate->segment  = (u16) d->segment;
    gate->bits     = d->bits;
    gate->offset_middle = (u16) (addr >> 16);
#ifndef CONFIG_X86_64
    gate->offset_high = (u32) (addr >> 32);
    gate->reserved  = 0;
#endif
}

low/middle/high      gate_desc    write_idt_entry      idt_table      native_write_id

```

```

static inline void native_write_idt_entry(gate_desc *idt, int entry, const gate_desc *gate)
{
    memcpy(&idt[entry], gate, sizeof(*gate));
}

entry          idt_table      system_vectors   bitmap        sys
vector         true    apic_ids

for_each_clear_bit_from(i, system_vectors, FIRST_SYSTEM_VECTOR) {
    entry = irq_entries_start + IDT_ALIGN * (i - FIRST_EXTERNAL_VECTOR);
    set_intr_gate(i, entry);
}

FIRST_SYSTEM_VECTOR  system_vectors  irq_entries_start  set_intr_gate  id
static __init void set_intr_gate(unsigned int n, const void *addr)
{
    struct idt_data data;

    init_idt_data(&data, n, addr);

    idt_setup_from_table(idt_table, &data, 1, false);
}

init_idt_data  idt_data       idt_setup_from_table  irq_entries_start  idt_table      false
i   system_vectors           idt_table             " "                  /           commo
>spurious_interrupt     CPU           " "                handle_spurious_interrupt
interrupt  /proc/interrupts          APIC  End-Of-Interrupt(EOI)
    spurious_entries_start

./arch/x86/include/asm/idtentry.h
SYM_CODE_START(spurious_entries_start)
    vector=FIRST_SYSTEM_VECTOR
    .rept NR_SYSTEM_VECTORS
    UNWIND_HINT_IRET_REGS
0 :
    ENDBR
    .byte 0x6a, vector
    jmp asm_spurious_interrupt
    /* Ensure that the above is IDT_ALIGN bytes max */
    .fill 0b + IDT_ALIGN - ., 1, 0xcc
    vector = vector+1
    .endr
SYM_CODE_END(spurious_entries_start)

asm_spurious_interrupt    asm_sysvec_reschedule_ipi  C

```

```

arch/x86/include/asm/idtentry.h

DECLARE_IDTENTRY_IRQ(X86_TRAP_OTHER, spurious_interrupt);

#define DECLARE_IDTENTRY_IRQ(vector, func)      \
DECLARE_IDTENTRY_ERRORCODE(vector, func)

#define DECLARE_IDTENTRY_ERRORCODE(vector, func)   \
asmlinkage void asm_##func(void);    \
asmlinkage void xen_asm_##func(void);    \
__visible void func(struct pt_regs *regs, unsigned long error_code)
.S

arch/x86/include/asm/idtentry.h

#define DECLARE_IDTENTRY_ERRORCODE(vector, func)   \
idtentry vector asm_##func func has_error_code=1

ipi
has_error_code 1      orig_ax
idt_map_in_cea

static void __init idt_map_in_cea(void)
{
/*
 * Set the IDT descriptor to a fixed read-only location in the cpu
 * entry area, so that the "sidt" instruction will not leak the
 * location of the kernel, and to defend the IDT against arbitrary
 * memory write vulnerabilities.
 */
cea_set_pte(CPU_ENTRY_AREA_RO_IDT_VADDR, __pa_symbol(idt_table),
            PAGE_KERNEL_RO);
idt_descr.address = CPU_ENTRY_AREA_RO_IDT;
}

CPU_ENTRY_AREA_RO_IDT_VADDR

arch/x86/include/asm/pgtable_areas.h
#define CPU_ENTRY_AREA_RO_IDT_VADDR ((void *)CPU_ENTRY_AREA_RO_IDT)

/* Single page reserved for the readonly IDT mapping: */
#define CPU_ENTRY_AREA_RO_IDT CPU_ENTRY_AREA_BASE

arch/x86/include/asm/pgtable_64_types.h
#define CPU_ENTRY_AREA_BASE (CPU_ENTRY_AREA_PGD << P4D_SHIFT)

#define CPU_ENTRY_AREA_PGD _AC(-4, UL)
#define P4D_SHIFT 39

```

```

CPU_ENTRY_AREA_RO_IDT_VADDR 0xfffffe0000000000          x86-
64      Documentation/arch/x86/x86_64/mm.rst 0000000000000000-00007fffffefff 128TB
ffffe8fffffffff 32TB  vmalloc/ioremap
CPU_ENTRY_AREA_RO_IDT_VADDR fffffe0000000000 fffffe0000000000-
fffffe7fffffff 512GB cpu_entry_area mapping      idt          Page-
Global           cr3 page global    tlb      cea_set_pte

void cea_set_pte(void *cea_vaddr, phys_addr_t pa, pgprot_t flags)
{
    unsigned long va = (unsigned long) cea_vaddr;
    pte_t pte = pfn_pte(pa >> PAGE_SHIFT, flags);

    /*
     * The cpu_entry_area is shared between the user and kernel
     * page tables. All of its ptes can safely be global.
     * _PAGE_GLOBAL gets reused to help indicate PROT_NONE for
     * non-present PTEs, so be careful not to set it in that
     * case to avoid confusion.
     */
    if (boot_cpu_has(X86_FEATURE_PGE) &&
        (pgprot_val(flags) & _PAGE_PRESENT))
        pte = pte_set_flags(pte, _PAGE_GLOBAL);

    set_pte_vaddr(va, pte);
}

idt_table      idt_table      idt_table      __pa_symbol
arch/x86/include/asm/page.h

#define __pa_symbol(x) \
__phys_addr_symbol(__phys_reloc_hide((unsigned long)(x)))

arch/x86/include/asm/page_64.h

#define __phys_addr_symbol(x) \
((unsigned long)(x) - __START_KERNEL_map + phys_base)

__START_KERNEL_map phys_base x86-64      __START_KERNEL_map
x86/include/asm/page_64_types.h

#define __START_KERNEL_map _AC(0xffffffff80000000, UL)
Documentation/arch/x86/x86_64/mm.rst
ffffffffff80000000 | -2 GB | ffffffff9fffffff | 512 MB | kernel text mapping, mapped to
idt_table      __phys_addr_symbol

```

```

/* Must be page-aligned because the real IDT is used in the cpu entry area */
static gate_desc idt_table[IDT_ENTRIES] __page_aligned_bss;

phys_base      __START_KERNEL_map    phys_base      phys_base __startup_64

/*
 * Compute the delta between the address I am compiled to run at
 * and the address I am actually running at.
 */
phys_base = load_delta = __START_KERNEL_map + p2v_offset;
__startup_64  startup_64

arch/x86/kernel/head_64.S

call __startup_64

startup_64 __startup_64           p2v_offset      phys_base   KASLR

p2v_offset = phys_base - __START_KERNEL_map
startup_64  rip  p2v_offset  __startup_64 C
phys_base = __START_KERNEL_map + p2v_offset
__START_KERNEL_map  phys_base  virt addr __START_KERNEL_map phys_base
                     virt addr/phys addr  __startup_64  [__ST
MMU CR3  phys_base  cpu  Page Fault
          cea_set_pte  pte  PAGE_KERNEL_RO

arch/x86/include/asm/pgtable_types.h

#define PAGE_KERNEL_RO __pgprot_mask(__PAGE_KERNEL_RO) | __ENC)
__ENC  AMD  __pgprot_mask

arch/x86/include/asm/pgtable_types.h

#define __pgprot_mask(x) __pgprot((x) & __default_kernel_pte_mask)
#define __pg(x)  __pgprot(x)
#define __pgprot(x) ((pgprot_t){(x)} )
typedef struct pgprot { pgprotval_t pgprot; } pgprot_t;

arch/x86/include/asm/pgtable_64_types.h

typedef unsigned long pgprotval_t;

__pgprot_mask     unsigned long     bit     __default_kernel_pte_mask     __PAGE_GLOBAL

```

```

arch/x86/mm/init.c: probe_page_size_mask

/* Except when with PTI where the kernel is mostly non-Global: */
if (cpu_feature_enabled(X86_FEATURE_PTI))
    __default_kernel_pte_mask &= ~_PAGE_GLOBAL;

arch/x86/mm/init_64.c

/* Bits allowed in normal kernel mappings: */
pteval_t __default_kernel_pte_mask __read_mostly = ~0;p
    __PAGE_KERNEL_RO

arch/x86/include/asm/pgtable_types.h

#define __PAGE_KERNEL_RO  (_PP|    0|    0|__A|__NX|    0|    0|__G)
__PP          __PP

arch/x86/include/asm/pgtable_types.h

#define __PP __PAGE_PRESENT
#define __PAGE_PRESENT (_AT(pteval_t, 1) << __PAGE_BIT_PRESENT)
#define __PAGE_BIT_PRESENT 0 /* is present */

```

bit                x86-64      4KB      intel sdm vol3      bit

Bit Position(s)	Contents
0 (P)	Present; must be 1 to map a 4-KByte page
1 (R/W)	Read/write; if 0, writes may not be allowed to the 4-KByte page referenced by this entry (see Section 4.6)
2 (U/S)	User/supervisor; if 0, user-mode accesses are not allowed to the 4-KByte page referenced by this entry (see Section 4.6)
3 (PWT)	Page-level write-through; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
4 (PD)	Page-level cache disable; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
5 (A)	Accessed; indicates whether software has accessed the 4-KByte page referenced by this entry (see Section 4.8)
6 (D)	Dirty; indicates whether software has written to the 4-KByte page referenced by this entry (see Section 4.8)
7 (PAT)	Indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
8 (G)	Global; if CR4.PGE = 1, determines whether the translation is global (see Section 4.10); ignored otherwise
11:9	Ignored
(M-1):12	Physical address of the 4-KByte page referenced by this entry
51:M	Reserved (must be 0)
58:52	Ignored
62:59	Protection key; if CR4.PKE = 1 or CR4.PKS = 1, this may control the page's access rights (see Section 4.6.2); otherwise, it is not used to control access rights
63 (X0)	If IA32_EFER.NXE = 1, execute-disable (if 1), instruction fetches are not allowed from the 4-KByte page controlled by this entry; see Section 4.8; otherwise, reserved (must be 0)

Figure 1: 4KB

```

cea_set_pte        cea_set_pte

void cea_set_pte(void *cea_vaddr, phys_addr_t pa, pgprot_t flags)
{
    unsigned long va = (unsigned long) cea_vaddr;
    pte_t pte = pfn_pte(pa >> PAGE_SHIFT, flags);

/*

```

```

* The cpu_entry_area is shared between the user and kernel
* page tables. All of its ptes can safely be global.
* _PAGE_GLOBAL gets reused to help indicate PROT_NONE for
* non-present PTEs, so be careful not to set it in that
* case to avoid confusion.
*/
if (boot_cpu_has(X86_FEATURE_PGE) &&
    (pgprot_val(flags) & _PAGE_PRESENT))
    pte = pte_set_flags(pte, _PAGE_GLOBAL);

    set_pte_vaddr(va, pte);
}

pte_t : 

typedef struct { pteval_t pte; } pte_t;
pteval_t 32/64 bit      64     unsigned long
typedef unsigned long pteval_t;

pte entry 64 bit pfn_pte mm          pte
static inline pte_t pfn_pte(unsigned long page_nr, pgprot_t pgprot)
{
    phys_addr_t pfn = (phys_addr_t)page_nr << PAGE_SHIFT;
    /* This bit combination is used to mark shadow stacks */
    WARN_ON_ONCE((pgprot_val(pgprot) & (_PAGE_DIRTY | _PAGE_RW)) ==
        _PAGE_DIRTY);
    pfn ^= protnone_mask(pgprot_val(pgprot));
    pfn &= PTE_PFN_MASK;
    return __pte(pfn | check_pgprot(pgprot));
}

page_nr      PAGE_SHIFT      pte          12      page_nr idt_table      PAGE_SH
      pfn_pte      pgprot      _PAGE_DIRTY      _PAGE_RW      " "      shadow
stacks      CET feature      (_PAGE_RW) CPU      _PAGE_DIRTY      (!_PAGE_RW)      _PAGE_
      _PAGE_DIRTY

/*
 * A clear pte value is special, and doesn't get inverted.
 *
 * Note that even users that only pass a pgprot_t (rather
 * than a full pte) won't trigger the special zero case,
 * because even PAGE_NONE has _PAGE_PROTNONE | _PAGE_ACSESSED
 * set. So the all zero case really is limited to just the
 * cleared page table entry case.
*/

```

```

static inline bool __pte_needs_invert(u64 val)
{
    return val && !(val & _PAGE_PRESENT);
}

/* Get a mask to xor with the page table entry to get the correct pfn. */
static inline u64 protnone_mask(u64 val)
{
    return __pte_needs_invert(val) ? ~0ull : 0;
}

x86 64      PTE      (PFN) PFN +      mprotect(PROT_NONE) "      Linux
      PROTNONE pfn      invert PFN " " "PROTNONE" val
!= 0 val & _PAGE_PRESENT!= 0      __pte_needs_invert(val) false pte PFN invert val
== 0 " " pte_clear      invert val!= 0 val & _PAGE_PRESENT
== 0 mprotect(PROT_NONE)      " " PFN      __pte_needs_invert(val) true invert PFN
pfn ^= protnone_mask(pgprot_val(pgprot));
PROTNONE      PFN      __PP  protnone_mask 0  0

/* Extracts the PFN from a (pte|pmd|pud|pgd)val_t of a 4KB page */
#define PTE_PFN_MASK ((pteval_t)PHYSICAL_PAGE_MASK)

        4KB PFN          64          PHYSICAL_PAGE_MASK

#define PHYSICAL_PAGE_MASK (((signed long)PAGE_MASK) & __PHYSICAL_MASK)
PAGE_MASK

#define PAGE_MASK (~(PAGE_SIZE - 1))
__PHYSICAL_MASK     bit
#define __PHYSICAL_MASK physical_mask

phys_addr_t physical_mask __ro_after_init = (1ULL << __PHYSICAL_MASK_SHIFT) - 1;
#define __PHYSICAL_MASK_SHIFT 52

__pte
#define __pte(x) native_make_pte(x)

static inline pte_t native_make_pte(pteval_t val)
{
    return (pte_t) { .pte = val };
}

        pfn_pte      cea_set_pte      present=1      global      present=0      global      Linux " "
present  global      PROT_NONE      present      global      __PAGE_GLOBAL      CR3      flu
pte_set_flags

```

```

static inline pte_t pte_set_flags(pte_t pte, pteval_t set)
{
    pteval_t v = native_pte_val(pte);

    return native_make_pte(v | set);
}

static inline pte_t native_make_pte(pteval_t val)
{
    return (pte_t) { .pte = val };
}

va pte    set_pte_vaddr  pte
arch/x86/mm/init_64.c

void set_pte_vaddr(unsigned long vaddr, pte_t pteval)
{
    pgd_t *pgd;
    p4d_t *p4d_page;

    pr_debug("set_pte_vaddr %lx to %lx\n", vaddr, native_pte_val(pteval));

    pgd = pgd_offset_k(vaddr);
    if (pgd_none(*pgd)) {
        printk(KERN_ERR
              "PGD FIXMAP MISSING, it should be setup in head.S!\n");
        return;
    }

    p4d_page = p4d_offset(pgd, 0);
    set_pte_vaddr_p4d(p4d_page, vaddr, pteval);
}

set_pte_vaddr                         crash vtop           vtop
                                         0xfffffe0000000000

crash> vtop 0xfffffe0000000000
      VIRTUAL          PHYSICAL
      ffffffe0000000000  5db2d000

PGD DIRECTORY: ffffffff3422000
PAGE DIRECTORY: 23ffc5067
    PUD: 23ffc5000 => 1003b8067
    PMD: 1003b8000 => 1003b9067
    PTE: 1003b9000 => 800000005db2d121
    PAGE: 5db2d000

```

```

PTE          PHYSICAL  FLAGS
800000005db2d121 5db2d000 (PRESENT|ACCESSED|GLOBAL|NX)

PAGE          PHYSICAL      MAPPING      INDEX CNT FLAGS
fffffde8e4176cb40 5db2d000           0         0 1 fffffc0002000 reserved

pgd_offset_k    init_mm      address pgd

include/linux/pgtable.h

/*
 * a shortcut which implies the use of the kernel's pgd, instead
 * of a process's
 */
#define pgd_offset_k(address) pgd_offset(&init_mm, (address))
#define pgd_offset(mm, address) pgd_offset_pgd((mm)->pgd, (address))
static inline pgd_t *pgd_offset_pgd(pgd_t *pgd, unsigned long address)
{
    return (pgd + pgd_index(address));
}
#define pgd_index(a) (((a) >> PGDIR_SHIFT) & (PTRS_PER_PGD - 1))

arch/x86/include/asm/pgtable_64_types.h

/*
 * PGDIR_SHIFT determines what a top-level page table entry can map
 */
#define PGDIR_SHIFT pgdir_shift
#define PTRS_PER_PGD 512

pgdir_shift    39

arch/x86/kernel/head64.c

unsigned int pgdir_shift __ro_after_init = 39;

      5 pgdir_shift      crash p pgdir_shift      39

crash> p pgdir_shift
pgdir_shift = $29 = 39

      pgd_offset_k    pgd_offset_k

      init_mm.pgd + (address >> 39 & 511) * 8

      init_mm.pgd

mm/init-mm.c

```

```

struct mm_struct init_mm = {
    ...
    .pgd  = swapper_pg_dir,
    ...
};

arch/x86/include/asm/pgtable_64.h

#define swapper_pg_dir init_top_pgt
init_top_pgt          PGD
arch/x86/kernel/head_64.S

SYM_DATA_START_PTI_ALIGNED(init_top_pgt)
    .quad    level3_ident_pgt - __START_KERNEL_map + _KERNPG_TABLE_NOENC
    .org     init_top_pgt + L4_PAGE_OFFSET*8, 0
    .quad    level3_ident_pgt - __START_KERNEL_map + _KERNPG_TABLE_NOENC
    .org     init_top_pgt + L4_START_KERNEL*8, 0
/* (2^48-(2*1024*1024*1024))/(2^39) = 511 */
    .quad    level3_kernel_pgt - __START_KERNEL_map + _PAGE_TABLE_NOENC
    .fill   PTI_USER_PGD_FILL,8,0
SYM_DATA_END(init_top_pgt)

.cpu          cpu
.crash      pgd init_top_pgt
                .quad      level3_ident_pg/level3_kernel_pgt
crash> p init_mm.pgd
$30 = (pgd_t *) 0xfffffffffb3422000 <init_top_pgt>
     8   1   8
0xfffffffffb3422000 + (0xfffffe0000000000 >> 39 & 511) * 8 = 0xfffffffffb3422fe0
crash rd      crash      23ffc5067
crash> rd 0xfffffffffb3422fe0
fffffffffffb3422fe0:  000000023ffc5067           gP.?.....
    pgd_offset_k          pgd          head_64.S
    p4d_offset

arch/x86/include/asm/pgtable.h

/* to find an entry in a page-table-directory. */
static inline p4d_t *p4d_offset(pgd_t *pgd, unsigned long address)
{
    if (!pgtable_l5_enabled())
        return (p4d_t *)pgd;
    return (p4d_t *)pgd_page_vaddr(*pgd) + p4d_index(address);
}

```

```

AMD      5  cat /proc/self/maps 12 0 4bit 48bit

arch/x86/include/asm/pgtable.h

static inline unsigned long pgd_page_vaddr(pgd_t pgd)
{
    return (unsigned long)_va((unsigned long)pgd_val(pgd) & PTE_PFN_MASK);
}

arch/x86/include/asm/page.h

#define __va(x) ((void *)((unsigned long)(x)+PAGE_OFFSET))

arch/x86/include/asm/page_types.h

#define PAGE_OFFSET ((unsigned long)__PAGE_OFFSET)

arch/x86/include/asm/page_64_types.h

#define __PAGE_OFFSET           page_offset_base
page_offset_base arch/x86/mm/kaslr.c      PAGE_OFFSET      crash  page_offset_base
crash> p/x page_offset_base
$38 = 0xfffff8a92c0000000

*pgd  pgd          p4d          PTE_PFN_MASK  PTE_PFN_MASK
p4d_page = 0xfffff8a92c000000+0x23ffc5000 = 0xfffff8a94fffc5000
crash rd      p4d pud

crash> rd 0xfffff8a94fffc5000
fffff8a94fffc5000: 00000001003b8067           g.;.....
crash

PUD: 23ffc5000 => 1003b8067

set_pte_vaddr_p4d
arch/x86/mm/init_64.c

void set_pte_vaddr_p4d(p4d_t *p4d_page, unsigned long vaddr, pte_t new_pte)
{
    p4d_t *p4d = p4d_page + p4d_index(vaddr);
    pud_t *pud = fill_pud(p4d, vaddr);

    __set_pte_vaddr(pud, vaddr, new_pte);
}

```

```

p4d_index

arch/x86/include/asm/pgtable.h

static inline unsigned long p4d_index(unsigned long address)
{
    return (address >> P4D_SHIFT) & (PTRS_PER_P4D - 1);
}

arch/x86/include/asm/pgtable_64_types.h

#define PTRS_PER_P4D  ptrs_per_p4d

arch/x86/boot/compressed/pgtable_64.c

unsigned int __section(".data") ptrs_per_p4d = 1;

crash      ptrs_per_p4d  1  p4d_index  0  p4d      p4d p4d_page  0xffff8a94ffc5000 fill_pud

arch/x86/mm/init_64.c

static pud_t *fill_pud(p4d_t *p4d, unsigned long vaddr)
{
    if (p4d_none(*p4d)) {
        pud_t *pud = (pud_t *)spp_getpage();
        p4d_populate(&init_mm, p4d, pud);
        if (pud != pud_offset(p4d, 0))
            printk(KERN_ERR "PAGETABLE BUG #01! %p <-> %p\n",
                  pud, pud_offset(p4d, 0));
    }
    return pud_offset(p4d, vaddr);
}

p4d          p4d          flag          spp_getpage
pud

/*
 * NOTE: This function is marked __ref because it calls __init function
 * (alloc_bootmem_pages). It's safe to do it ONLY when after_bootmem == 0.
 */
static __ref void *spp_getpage(void)
{
    void *ptr;

    if (after_bootmem)
        ptr = (void *) get_zeroed_page(GFP_ATOMIC);
    else
        ptr = memblock_alloc(PAGE_SIZE, PAGE_SIZE);

```

```

if (!ptr || ((unsigned long)ptr & ~PAGE_MASK)) {
    panic("set_pte_phys: cannot allocate page data %s\n",
          after_bootmem ? "after bootmem" : "");
}

pr_debug("spp_getpage %p\n", ptr);

return ptr;
}

__ref           init/exit      /           init/exit  init/exit   modpost

/*
 * modpost check for section mismatches during the kernel build.
 * A section mismatch happens when there are references from a
 * code or data section to an init section (both code or data).
 * The init sections are (for most archs) discarded by the kernel
 * when early init has completed so all such references are potential bugs.
 * For exit sections the same issue exists.
 *
 * The following markers are used for the cases where the reference to
 * the __init / __exit section (code or data) is valid and will teach
 * modpost not to issue a warning. Intended semantics is that a code or
 * data tagged __ref can reference code or data from init section without
 * producing a warning (of course, no warning does not mean code is
 * correct, so optimally document why the __ref is needed and why it's OK).
 *
 * The markers follow same syntax rules as __init / __initdata.
 */
#define __ref           __section(".ref.text") noinline

spp_getpage

/*
 * NOTE: This function is marked __ref because it calls __init function
 * (alloc_bootmem_pages). It's safe to do it ONLY when after_bootmem == 0.
 */
static __ref void *spp_getpage(void)
{
    void *ptr;

    if (after_bootmem)
        ptr = (void *) get_zeroed_page(GFP_ATOMIC);
    else
        ptr = memblock_alloc(PAGE_SIZE, PAGE_SIZE);
}

```

```

if (!ptr || ((unsigned long)ptr & ~PAGE_MASK)) {
    panic("set_pte_phys: cannot allocate page data %s\n",
          after_bootmem ? "after bootmem" : "");
}

pr_debug("spp_getpage %p\n", ptr);

return ptr;
}

memblock    membblock_alloc      buddy system  get_zeroed_page   p4d_populate   p4d
arch/x86/include/asm/pgalloc.h

static inline void p4d_populate(struct mm_struct *mm, p4d_t *p4d, pud_t *pud)
{
    paravirt_alloc_pud(mm, __pa(pud) >> PAGE_SHIFT);
    set_p4d(p4d, __p4d(_PAGE_TABLE | __pa(pud)));
}

_PAGE_TABLE      pte
/*
 * Page tables needs to have Write=1 in order for any lower PTEs to be
 * writable. This includes shadow stack memory (Write=0, Dirty=1)
 */
#define _PAGE_TABLE  (_PP|_RW|_USR|__A| 0|__D| 0| 0| _ENC)
page table      __pa
arch/x86/include/asm/page.h

#ifndef __pa
#define __pa(x)  __phys_addr((unsigned long)(x))
#endif

arch/x86/include/asm/page_64.h

#define __phys_addr(x)  __phys_addr_nodebug(x)

static __always_inline unsigned long __phys_addr_nodebug(unsigned long x)
{
    unsigned long y = x - __START_KERNEL_map;

    /* use the carry flag to determine if x was < __START_KERNEL_map */
    x = y + ((x > y) ? phys_base : (__START_KERNEL_map - PAGE_OFFSET));

    return x;
}

```

```

    __phys_addr_nodebug      __phys_addr_symbol          x > __START_KERNEL_map
= x - __START_KERNEL_map + phy_base x < __START_KERNEL_map y     x
> y      phy_addr(x) = x - __START_KERNEL_map + __START_KERNEL_map
- PAGE_OFFSET = x - PAGE_OFFSET           PAGE_OFFSET      PAGE_OFFSET
    __p4d      flag _PAGE_TABLE

arch/x86/include/asm/pgtable.h

#ifndef set_p4d
#define set_p4d(p4dp, p4d) native_set_p4d(p4dp, p4d)
#endif

static inline void native_set_p4d(p4d_t *p4dp, p4d_t p4d)
{
    pgd_t pgd;

    if (pgtable_l5_enabled() ||
        !IS_ENABLED(CONFIG_MITIGATION_PAGE_TABLE_ISOLATION)) {
        WRITE_ONCE(*p4dp, p4d);
        return;
    }

    pgd = native_make_pgd(native_p4d_val(p4d));
    pgd = pti_set_user_pttbl((pgd_t *)p4dp, pgd);
    WRITE_ONCE(*p4dp, native_make_p4d(native_pgd_val(pgd)));
}

WRITE_ONCE    pti /           p4d p4dp      WRITE_ONCE    *p4dp           p4dp
    __set_pte_vaddr           spp_getpage           CPU_ENTRY_AREA_RO_IDT_V
    cea_set_pte      /       crash         vtop 0xfffffe0000000000

crash> vtop 0xfffffe0000000000
VIRTUAL          PHYSICAL
fffffe0000000000 5db2d000

PGD DIRECTORY: ffffffff3422000
PAGE DIRECTORY: 23ffc5067
    PUD: 23ffc5000 => 1003b8067
    PMD: 1003b8000 => 1003b9067
    PTE: 1003b9000 => 800000005db2d121
PAGE: 5db2d000

    PTE          PHYSICAL  FLAGS
800000005db2d121 5db2d000 (PRESENT|ACCESSED|GLOBAL|NX)

    PAGE          PHYSICAL      MAPPING      INDEX CNT FLAGS
fffffde8e4176cb40 5db2d000          0          0 1 ffffffc0002000 reserved

```

```

5db2d000      idt_table           idt_table      pte          crash    idt_ta
crash> p &idt_table
$71 = (gate_desc (*)[256]) 0xfffffffffb3f2d000 <idt_table>
crash> vtop 0xfffffffffb3f2d000
VIRTUAL          PHYSICAL
fffffffffb3f2d000 5db2d000

PGD DIRECTORY: ffffffb3422000
PAGE DIRECTORY: 5d027067
  PUD: 5d027ff0 => 5d028063
  PMD: 5d028cf8 => 1003ba063
  PTE: 1003ba968 => 800000005db2d121
PAGE: 5db2d000

      PTE          PHYSICAL  FLAGS
800000005db2d121 5db2d000 (PRESENT|ACCESSED|GLOBAL|NX)
      PAGE          PHYSICAL      MAPPING      INDEX CNT FLAGS
fffffde8e4176cb40 5db2d000          0          0 1 ffffffc0002000 reserve

      idt_table      idt_table      CPU_ENTRY_AREA_RO_IDT_VADDR 0xfffffe0000000000
      idt  idt_table      CPU_ENTRY_AREA_RO_IDT_VADDR          CPU_
only   crash     1   x86      _PAGE_RW=0
      load_idt  0xfffffe0000000000 idtr

load_idt(&idt_descr);

#define load_idt(dtr)    native_load_idt(dtr)
static __always_inline void native_load_idt(const struct desc_ptr *dtr)
{
    asm volatile("lidt %0::"m" (*dtr));
}

idt_setup_apic_and_irq_gates  set_memory_ro->change_page_attr_clear idt_table _PAGE_RW _PAGE_RW

/* Make the IDT table read only */
set_memory_ro((unsigned long)&idt_table, 1);

int set_memory_ro(unsigned long addr, int numpages)
{
    return change_page_attr_clear(&addr, numpages, __pgprot(_PAGE_RW | _PAGE_DIRTY), 0);
}

```

x86-64

## 2

```
          irq_entries_start           SYM_CODE_START      push    jmp  asm_common_
vector
          asm_common_interrupt        "       C   "
arch/x86/include/asm/idtentry.h

/* Device interrupts common/spurious */
DECLARE_IDTENTRY_IRQ(X86_TRAP_OTHER, common_interrupt);

          DECLARE_IDTENTRY_IRQ .c .S
arch/x86/include/asm/idtentry.h

#ifndef __ASSEMBLER__
<**
 * DECLARE_IDTENTRY_ERRORCODE - Declare functions for simple IDT entry points
 *     Error code pushed by hardware
 * @vector: Vector number (ignored for C)
 * @func: Function name of the entry point
 *
 * Declares three functions:
 * - The ASM entry point: asm_##func
 * - The XEN PV trap entry point: xen_##func (maybe unused)
 * - The C handler called from the ASM entry point
 *
 * Same as DECLARE_IDTENTRY, but has an extra error_code argument for the
 * C-handler.
*/
#define DECLARE_IDTENTRY_ERRORCODE(vector, func) \
asmlinkage void asm_##func(void); \
asmlinkage void xen_asm_##func(void); \
__visible void func(struct pt_regs *regs, unsigned long error_code)

<**
 * DECLARE_IDTENTRY_IRQ - Declare functions for device interrupt IDT entry
 *     points (common/spurious)
 * @vector: Vector number (ignored for C)
 * @func: Function name of the entry point
 *
 * Maps to DECLARE_IDTENTRY_ERRORCODE()
*/
#define DECLARE_IDTENTRY_IRQ(vector, func) \
DECLARE_IDTENTRY_ERRORCODE(vector, func)
```

```

#else /* !_ASSEMBLER_ */
/* Entries for common/spurious (device) interrupts */
#define DECLARE_IDTENTRY_IRQ(vector, func) \
    idtentry_irq vector func

c     asm_common_interrupt           .S      idtentry_irq      idtentry.h  :
./arch/x86/kernel/idt.c:14:#include <asm/idtentry.h>
.c
./arch/x86/entry/entry_64.S:552:#include <asm/idtentry.h>

          Linux "      C "      /      stub CPU
-> C handler-> ->iretq C handler      +C
asm_common_interrupt

ffffffffff81e00230:    f3 0f 1e fa      endbr64
ffffffffff81e00234:    6a 20      push $0x20
ffffffffff81e00236:    e9 05 13 00 00      jmp ffffffff81e01540 <asm_common_interrupt
ffffffffff81e0023b:    cc      int3
ffffffffff81e0023c:    cc      int3
ffffffffff81e0023d:    cc      int3
ffffffffff81e0023e:    cc      int3
ffffffffff81e0023f:    cc      int3

asm_common_interrupt      arch/x86/entry/entry_64.S X86-64   X86-
32 arch/x86/entry/entry_32.S

/* Device interrupts common/spurious */
DECLARE_IDTENTRY_IRQ(X86_TRAP_OTHER, common_interrupt);

/*
 * Dummy trap number so the low level ASM macro vector number checks do not
 * match which results in emitting plain IDTENTRY stubs without bells and
 * whistles.
 */
#define X86_TRAP_OTHER 0xFFFF

X86_TRAP_OTHER      vector number  irq_entries_start push

#define DECLARE_IDTENTRY_IRQ(vector, func) \
    idtentry_irq vector func

idtentry_irq entry_64.S
arch/x86/entry/entry_64.S

/*
 * Interrupt entry/exit.

```

```

*
* The interrupt stubs push (vector) onto the stack, which is the error_code
* position of idtentry exceptions, and jump to one of the two idtentry points
* (common/spurious).
*
* common_interrupt is a hotpath, align it to a cache line
*/
.macro idtentry_irq vector cfunc
.p2align CONFIG_X86_L1_CACHE_SHIFT
idtentry \vector asm_\cfunc \cfunc has_error_code=1
.endm

.p2align      2^CONFIG_X86_L1_CACHE_SHIFT  CONFIG_X86_L1_CACHE_SHIFT 6 64
L1 cache line          cache line  cache miss
        asm_common_interrupt common_interrupt idtentry    asm_common_interrupt common_in
arch/x86/entry/entry_64.S

/***
* idtentry - Macro to generate entry stubs for simple IDT entries
* @vector: Vector number
* @asmsym: ASM symbol for the entry point
* @cfunc: C function to be called
* @has_error_code: Hardware pushed error code on stack
*
* The macro emits code to set up the kernel context for straight forward
* and simple IDT entries. No IST stack, no paranoid entry checks.
*/
.macro idtentry vector asmsym cfunc has_error_code:req
SYM_CODE_START(\asmsym)

.if \vector == X86_TRAP_BP
/* #BP advances %rip to the next instruction */
UNWIND_HINT_IRET_ENTRY offset=\has_error_code*8 signal=0
.else
UNWIND_HINT_IRET_ENTRY offset=\has_error_code*8
.endif

ENDBR
ASM_CLAC
cld

.if \has_error_code == 0
pushq $-1 /* ORIG_RAX: no syscall to restart */
.endif

.if \vector == X86_TRAP_BP

```

```

/*
 * If coming from kernel space, create a 6-word gap to allow the
 * int3 handler to emulate a call instruction.
 */
testb $3, CS-ORIG_RAX(%rsp)
jnz .Lfrom_usermode_no_gap_\
.rept 6
pushq 5*8(%rsp)
.endr
UNWIND_HINT_IRET_REGS offset=8
.Lfrom_usermode_no_gap_@:
.endif

idtentry_body \cfunc \has_error_code

_ASM_NOKPROBE(\asmsym)
SYM_CODE_END(\asmsym)
.endm

asm_common_interrupt

fffffffff81e01540 <asm_common_interrupt>:
fffffffff81e01540: f3 0f 1e fa      endbr64
fffffffff81e01544: 90               nop
fffffffff81e01545: 90               nop
fffffffff81e01546: 90               nop
fffffffff81e01547: fc               cld
fffffffff81e01548: e8 e3 06 00 00   call   ffffffff81e01c30 <error_entry>
fffffffff81e0154d: 48 89 c4       mov    %rax,%rsp
fffffffff81e01550: 48 89 e7       mov    %rsp,%rdi
fffffffff81e01553: 48 b8 74 24 78   mov    0x78(%rsp),%rsi
fffffffff81e01558: 48 c7 44 24 78   movq   $0xfffffffffffffff,0x78(%rsp)
fffffffff81e0155f: ff ff
fffffffff81e01561: e8 5a 0c ef ff   call   ffffffff81cf21c0 <common_interrupt>
fffffffff81e01566: e9 05 08 00 00   jmp    ffffffff81e01d70 <error_return>
fffffffff81e0156b: 66 66 2e 0f 1f 84 00  data16 cs nopw 0x0(%rax,%rax,1)
fffffffff81e01572: 00 00 00 00
fffffffff81e01576: 66 2e 0f 1f 84 00 00  cs nopw 0x0(%rax,%rax,1)
fffffffff81e0157d: 00 00 00

          SYM_CODE_START  .global
if-else     IDT      ORC unwinder      #BP breakpoint        RIP
ENDBR      ASM_CLAC  ASM_CLAC        .S .c             .S      .c

arch/x86/include/asm/smap.h

#endif __ASSEMBLER__

```

```

#define ASM_CLAC \
    ALTERNATIVE "", "clc", X86_FEATURE_SMAP
arch/x86/include/asm/alternative.h

/*
 * Issue one struct alt_instr descriptor entry (need to put it into
 * the section .altinstructions, see below). This entry contains
 * enough information for the alternatives patching code to patch an
 * instruction. See apply_alternatives().
 */
.macro altinstr_entry orig alt ft_flags orig_len alt_len
    .long \orig - .
    .long \alt - .
    .4byte \ft_flags
    .byte \orig_len
    .byte \alt_len
.endm

/*
 * Define an alternative between two instructions. If @feature is
 * present, early code in apply_alternatives() replaces @oldinstr with
 * @newinstr. ".skip" directive takes care of proper instruction padding
 * in case @newinstr is longer than @oldinstr.
 */
#define __ALTERNATIVE(oldinst, newinst, flag) \
740:          \
    oldinst ;      \
741:          \
    .skip -(((744f-743f)-(741b-740b)) > 0) * ((744f-743f)-(741b-740b)),0x90 ; \
742:          \
    .pushsection .altinstructions,"a" ;      \
    altinstr_entry 740b,743f,flag,742b-740b,744f-743f ; \
    .popsection ;      \
    .pushsection .altinstr_replacement,"ax" ;      \
743:          \
    newinst ;      \
744:          \
    .popsection ;      \

.macro ALTERNATIVE oldinstr, newinstr, ft_flags
    __ALTERNATIVE(\oldinstr, \newinstr, \ft_flags)
.endm

__ALTERNATIVE    740      (oldinst)      741      skip      newinst oldinst  nop(0x90)
743f) 743-744      (741b-740b) 741-742      b f      b skip      backward f skip  skip  forward

```

```

feature flag      +          apply_alternatives CPU .altinstrs
.altinstr_replacement " " " "    nop     patch
.asm_common_interrupt X86_FEATURE_SMAP feature
Check           CR4   SMAP(Supervisor Mode Access Prevention)
                cld   (DF, Direction Flag) DF EFLAGS DF=0
MOVS, CMPS, STOS, LODS, SCAS DF = 1
                error code stub push vector syscall number has_error_code orig_ax
-1             error code intel sdm vol3 6.3.1 error code vector 14 Page
Fault error code vector 0 Divide Error error code intel sdm vol3 6.13 error
code   pushq -1      push error code / push procedure handler
-1     handler      intel sdm vol3 6.12.1

```

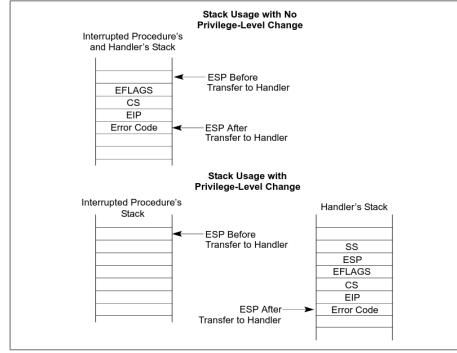


Figure 2: /

```

int3      6 * 8      call      int3      call  testb  cs      cs  rsp
+ 16    CS-ORIG_RAX offset      16

#define ORIG_RAX 120
#define CS 136

pt_regs      orig_rax error code      syscall number pr_regs      pt_regs
-1

/*
 * C ABI says these regs are callee-preserved. They aren't saved on kernel entry
 * unless syscall needs a complete, fully filled "struct pt_regs".
 */
#define R15 0
#define R14 8
#define R13 16
#define R12 24
#define RBP 32
#define RBX 40
/* These regs are callee-clobbered. Always saved on kernel entry. */
#define R11 48

```

```

#define R10 56
#define R9 64
#define R8 72
#define RAX 80
#define RCX 88
#define RDX 96
#define RSI 104
#define RDI 112
/*
 * On syscall entry, this is syscall#. On CPU exception, this is error code.
 * On hw interrupt, it's IRQ number:
 */
#define ORIG_RAX 120
/* Return frame for iretq */
#define RIP 128
#define CS 136
#define EFLAGS 144
#define RSP 152
#define SS 160
#endif /* __ASSEMBLER__ */

/* top of stack page */
#define FRAME_SIZE 168

struct pt_regs {
/*
 * C ABI says these regs are callee-preserved. They aren't saved on
 * kernel entry unless syscall needs a complete, fully filled
 * "struct pt_regs".
 */
unsigned long r15;
unsigned long r14;
unsigned long r13;
unsigned long r12;
unsigned long bp;
unsigned long bx;

/* These regs are callee-clobbered. Always saved on kernel entry. */
unsigned long r11;
unsigned long r10;
unsigned long r9;
unsigned long r8;
unsigned long ax;
unsigned long cx;
unsigned long dx;
unsigned long si;

```

```

unsigned long di;

/*
 * orig_ax is used on entry for:
 * - the syscall number (syscall, sysenter, int80)
 * - error_code stored by the CPU on traps and exceptions
 * - the interrupt number for device interrupts
 *
 * A FRED stack frame starts here:
 *   1) It always includes an error code;
 *
 *   2) The return frame for ERET[US] starts here, but
 *      the content of orig_ax is ignored.
 */
unsigned long orig_ax;

/* The IRETQ return frame starts here */
unsigned long ip;

union {
    /* CS selector */
    u16 cs;
    /* The extended 64-bit data slot containing CS */
    u64 csx;
    /* The FRED CS extension */
    struct fred_cs fred_cs;
}.*;

unsigned long flags;
unsigned long sp;

union {
    /* SS selector */
    u16 ss;
    /* The extended 64-bit data slot containing SS */
    u64 ssx;
    /* The FRED SS extension */
    struct fred_ss fred_ss;
}.*;

/*
 * Top of stack on IDT systems, while FRED systems have extra fields
 * defined above for storing exception related information, e.g. CR2 or
 * DR6.
 */
}.*;

```

```

orig_ax      orig_ax      syscall number      error code      intel
sdm vol3  error code   6.13 Table 6-1      orig_ax      vector      error
code       irq_entries_start push      orig_ax      error
code       identry    cfunc common_interrupt has_error_code      identry_body

/***
 * idtentry_body - Macro to emit code calling the C function
 * @cfunc:  C function to be called
 * @has_error_code: Hardware pushed error code on stack
 */
.macro idtentry_body cfunc has_error_code:req

/*
 * Call error_entry() and switch to the task stack if from userspace.
 *
 * When in XENPV, it is already in the task stack, and it can't fault
 * for native_iret() nor native_load_gs_index() since XENPV uses its
 * own pvpns for IRET and load_gs_index(). And it doesn't need to
 * switch the CR3. So it can skip invoking error_entry().
 */
ALTERNATIVE "call error_entry; movq %rax, %rsp", \
            "call xen_error_entry", X86 FEATURE_XENPV

ENCODE_FRAME_POINTER
UNWIND_HINT_REGS

movq %rsp, %rdi /* pt_regs pointer into 1st argument*/

.if \has_error_code == 1
    movq ORIG_RAX(%rsp), %rsi /* get error code into 2nd argument*/
    movq $-1, ORIG_RAX(%rsp) /* no syscall to restart */
.endif

/* For some configurations \cfunc ends up being a noreturn. */
ANNOTATE_REACHABLE
call \cfunc

jmp error_return
.endm

ALTERNATIVE Xen paravirtual guest           call error_entry error_entry
arch/x86/entry/entry_64.S

SYM_CODE_START(error_entry)
ANNOTATE_NOENDBR
UNWIND_HINT_FUNC

```

```

PUSH_AND_CLEAR_REGS save_ret=1
ENCODE_FRAME_POINTER 8

testb $3, CS+8(%rsp)
jz .Lerror_kernelspace

/*
 * We entered from user mode or we're pretending to have entered
 * from user mode due to an IRET fault.
 */
swapgs
FENCE_SWAPGS_USER_ENTRY
/* We have user CR3. Change to kernel CR3. */
SWITCH_TO_KERNEL_CR3 scratch_reg=%rax
IBRS_ENTER
UNTRAIN_RET_FROM_CALL

leaq 8(%rsp), %rdi /* arg0 = pt_regs pointer */
/* Put us onto the real thread stack. */
jmp sync_regs

/*
 * There are two places in the kernel that can potentially fault with
 * usergs. Handle them here. B stepping K8s sometimes report a
 * truncated RIP for IRET exceptions returning to compat mode. Check
 * for these here too.
 */
.Lerror_kernelspace:
leaq native_irq_return_iret(%rip), %rcx
cmpq %rcx, RIP+8(%rsp)
je .Lerror_bad_iret
movl %ecx, %eax /* zero extend */
cmpq %rax, RIP+8(%rsp)
je .Lbstep_iret
cmpq $.Lgs_change, RIP+8(%rsp)
jne .Lerror_entry_done_lfence

/*
 * hack: .Lgs_change can fail with user gsbase. If this happens, fix up
 * gsbase and proceed. We'll fix up the exception and land in
 * .Lgs_change's error handler with kernel gsbase.
 */
swapgs

/*

```

```

 * Issue an LFENCE to prevent GS speculation, regardless of whether it is a
 * kernel or user gsbase.
 */
.Lerror_entry_done_lfence:
FENCE_SWAPGS_KERNEL_ENTRY
CALL_DEPTH_ACCOUNT
leaq 8(%rsp), %rax /* return pt_regs pointer */
VALIDATE_UNRET_END
RET

.Lbstep_iret:
/* Fix truncated RIP */
movq %rcx, RIP+8(%rsp)
/* fall through */

.Lerror_bad_iret:
/*
 * We came from an IRET to user mode, so we have user
 * gsbase and CR3. Switch to kernel gsbase and CR3:
 */
swapgs
FENCE_SWAPGS_USER_ENTRY
SWITCH_TO_KERNEL_CR3 scratch_reg=%rax
IBRS_ENTER
UNTRAIN_RET_FROM_CALL

/*
 * Pretend that the exception came from user mode: set up pt_regs
 * as if we faulted immediately after IRET.
 */
leaq 8(%rsp), %rdi /* arg0 = pt_regs pointer */
call fixup_bad_iret
mov %rax, %rdi
jmp sync_regs
SYM_CODE_END(error_entry)

PUSH_AND_CLEAR_REGS
arch/x86/entry/calling.h

.macro PUSH_AND_CLEAR_REGS rdx=%rdx rcx=%rcx rax=%rax save_ret=0 clear_bp=1 unwind_hint=1
PUSH_REGS rdx=\rdx, rcx=\rcx, rax=\rax, save_ret=\save_ret unwind_hint=\unwind_hint
CLEAR_REGS clear_bp=\clear_bp

/*
 * 64-bit system call stack frame layout defines and helpers,
 * for assembly code:

```

```

*/
.macro PUSH_REGS rdx=%rdx rcx=%rcx rax=%rax save_ret=0 unwind_hint=1
.if \save_ret
pushq %rsi /* pt_regs->si */
movq 8(%rsp), %rsi /* temporarily store the return address in %rsi */
movq %rdi, 8(%rsp) /* pt_regs->di (overwriting original return address) */
/* We just clobbered the return address - use the IRET frame for unwinding: */
UNWIND_HINT_IRET_REGS offset=3*8
.else
pushq %rdi /* pt_regs->di */
pushq %rsi /* pt_regs->si */
.endif
pushq \rdx /* pt_regs->dx */
pushq \rcx /* pt_regs->cx */
pushq \rax /* pt_regs->ax */
pushq %r8 /* pt_regs->r8 */
pushq %r9 /* pt_regs->r9 */
pushq %r10 /* pt_regs->r10 */
pushq %r11 /* pt_regs->r11 */
pushq %rbx /* pt_regs->rbx */
pushq %rbp /* pt_regs->rbp */
pushq %r12 /* pt_regs->r12 */
pushq %r13 /* pt_regs->r13 */
pushq %r14 /* pt_regs->r14 */
pushq %r15 /* pt_regs->r15 */

.if \unwind_hint
UNWIND_HINT_REGS
.endif

.if \save_ret
pushq %rsi /* return address on top of stack */
.endif
.endm

PUSH_REGS pushq          error_entry      15      rsp
                   rsi   rsi   rsp+8  push rsi   error_entry    mov   rsp   error_entry   rsi
rsi rsp     8   rsi   rsi       mov rdi   error_entry   error_entry   rsi   rsp+8
                   CLEAR_REGS 15

.macro CLEAR_REGS clear_bp=1
/*
 * Sanitize registers of values that a speculation attack might
 * otherwise want to exploit. The lower registers are likely clobbered
 * well before they could be put to use in a speculative execution

```

```

* gadget.
*/
xorl %esi, %esi /* nospec si */
xorl %edx, %edx /* nospec dx */
xorl %ecx, %ecx /* nospec cx */
xorl %r8d, %r8d /* nospec r8 */
xorl %r9d, %r9d /* nospec r9 */
xorl %r10d, %r10d /* nospec r10 */
xorl %r11d, %r11d /* nospec r11 */
xorl %ebx, %ebx /* nospec rbx */
.if \clear_bp
xorl %ebp, %ebp /* nospec rbp */
.endif
xorl %r12d, %r12d /* nospec r12 */
xorl %r13d, %r13d /* nospec r13 */
xorl %r14d, %r14d /* nospec r14 */
xorl %r15d, %r15d /* nospec r15 */

.endm

error_entry    CS+8      code segment register          testb 3 code
segment register AND      Zero Flags testb ZF = 1 code seg-
segment register 0 / cpu   jz error_kernelspace / cpu   code
segment register 11

swapgs gs     MSR C0000102H IA32_KERNEL_GS_BASE           x86-
64 gs.base     TLS     GS.base per-cpu
per-CPU swapgs     swapgs gs.base " " per-cpu
CPU         swapgs gs.base   tls
               FENCE_SWAPGS_USER_ENTRY

.macro FENCE_SWAPGS_USER_ENTRY
ALTERNATIVE "", "lfence", X86_FEATURE_FENCE_SWAPGS_USER
.endm

nop lfence 3 nop lfence      memory fence      load fence      store
A B" LFENCE CPU      Producer      flag=1 Consumer
flag=1 flag=1           LFENCE consumer flag=1 consumer LFENCE flag
cpu

ffffffffff81001c8d:       65 48 8b 15 1b 31 c5      mov      %gs:0x2c5311b(%rip),%rdx # fff
                           swapgs     mov     swapgs     SWITCH_TO_KERNEL_CR3

.macro SWITCH_TO_KERNEL_CR3 scratch_reg:req
ALTERNATIVE "jmp .Lend_\@", "", X86_FEATURE_PTI
mov %cr3, \scratch_reg
ADJUST_KERNEL_CR3 \scratch_reg
mov \scratch_reg, %cr3

```

```

.Lend_\@:
.endm

PTI /           CONFIG_MITIGATION_PAGE_TABLE_ISOLATION      ALTERNATIVE

.macro ADJUST_KERNEL_CR3 reg:req
ALTERNATIVE "", "SET_NOFLUSH_BIT \reg", X86_FEATURE_PCID
/* Clear PCID and "MITIGATION_PAGE_TABLE_ISOLATION bit", point CR3 at kernel pagetables: */
andq    $(~PTI_USER_PGTABLE_AND_PCID_MASK), \reg
.endm

.macro SET_NOFLUSH_BIT reg:req
bts $X86_CR3_PCID_NOFLUSH_BIT, \reg
.endm

#define X86_CR3_PCID_NOFLUSH_BIT 63 /* Preserve old PCID */

pcid   feature Intel intel sdm vol3 4.10.1      TLB          tlb      tlb      cr3
63 Intel "no flush" CR3.NOFLUSH mov cr3      tlb          intel sdm vol3
4.10.4.1          CR4.PCIDE      X86_FEATURE_PCID      feature
          CR4.PCIDE = 0      PCID = 000H  TLB          Global Pages      PCID
= 000H          paging-structure caches      /      PCID  CR4.PCIDE=0 CPU          PCID=0  CR3      PC
          CR4.PCIDE = 1      MOV CR3  63 0      12 bits 11:0  PCID      TL-
B          PCID          CPU      PCID      PCID  CR3          PCID  12          PCID      TLE
63=0  CR3  PCID TLB      PCID          TLB
          CR4.PCIDE = 1      63 1      TLB
          invlpg/invpcid  tlb          intel sdm vol3 4.10.4.1
          cpu      pcide      PTI feature X86_FEATURE_PTI  ADJUST_KERNEL_CR3      bts CR3
63 1  mov cr3  CR4.PCIDE = 1 CR3.bit63 = 1      TLB
          ADJUST_KERNEL_CR3  andq  PTI_USER_PGTABLE_AND_PCID_MASK

/*
 * MITIGATION_PAGE_TABLE_ISOLATION PGDs are 8k.  Flip bit 12 to switch between the two
 * halves:
 */
#define PTI_USER_PGTABLE_BIT PAGE_SHIFT
#define PTI_USER_PGTABLE_MASK (1 << PTI_USER_PGTABLE_BIT)
#define PTI_USER_PCID_BIT X86_CR3_PTI_PCID_USER_BIT
#define PTI_USER_PCID_MASK (1 << PTI_USER_PCID_BIT)
#define PTI_USER_PGTABLE_AND_PCID_MASK (PTI_USER_PCID_MASK | PTI_USER_PGTABLE_MASK)

#ifndef CONFIG_MITIGATION_PAGE_TABLE_ISOLATION
# define X86_CR3_PTI_PCID_USER_BIT 11
#endif

X86_CR3_PTI_PCID_USER_BIT      pcid      pcid 12      user/kernel pcid  1      pcid  0      pcid
2048          PCID          bit11=1 2048

```

```

PAGE_SHIFT 12 bit 12      4KB   8KB       8KB   MITIGATION_PAGE_TABLE_ISOLATION
11 bit 12   ADJUST_KERNEL_CR3      0xffffffffffffe7ff  bit 11 bit 12 0 cr3
pcid      asic      pcid      4096 pcid      6

#define TLB_NR_DYN_ASIDS 6

CPU Linux      6 mm_struct      ASID           task          6 struct
tlb_state per-cpu ASID      cache line

static inline u16 user_pcid(u16 asid)
{
    u16 ret = kern_pcid(asid);
#ifdef CONFIG_MITIGATION_PAGE_TABLE_ISOLATION
    ret |= 1 << X86_CR3_PTI_PCID_USER_BIT;
#endif
    return ret;
}

/*
 * Given @asid, compute kPCID
 */
static inline u16 kern_pcid(u16 asid)
{
    VM_WARN_ON_ONCE(asid > MAX_ASID_AVAILABLE);

#ifdef CONFIG_MITIGATION_PAGE_TABLE_ISOLATION
/*
 * Make sure that the dynamic ASID space does not conflict with the
 * bit we are using to switch between user and kernel ASIDs.
 */
BUILD_BUG_ON(TLB_NR_DYN_ASIDS >= (1 << X86_CR3_PTI_PCID_USER_BIT));

/*
 * The ASID being passed in here should have respected the
 * MAX_ASID_AVAILABLE and thus never have the switch bit set.
 */
VM_WARN_ON_ONCE(asid & (1 << X86_CR3_PTI_PCID_USER_BIT));
#endif
/*
 * The dynamically-assigned ASIDs that get passed in are small
 * (<TLB_NR_DYN_ASIDS). They never have the high switch bit set,
 * so do not bother to clear it.
 *
 * If PCID is on, ASID-aware code paths put the ASID+1 into the
 * PCID bits. This serves two purposes. It prevents a nasty
 * situation in which PCID-unaware code saves CR3, loads some other

```

```

        * value (with PCID == 0), and then restores CR3, thus corrupting
        * the TLB for ASID 0 if the saved ASID was nonzero. It also means
        * that any bugs involving loading a PCID-enabled CR3 with
        * CR4.PCIDE off will trigger deterministically.
        */
    return asid + 1;
}

kern_pcid(asid)    asid+1    ASID=0  user_pcid(asid) kern_pcid(asid)      PTI      PCID
flush   PTI
          pcid    asid+1  asid=0 pcid=0      pcid-unaware
          user_pcid  native_flush_tlb_one_user

STATIC_NOPV void native_flush_tlb_one_user(unsigned long addr)
{
    u32 loaded_mm_asid;
    bool cpu_pcide;

    /* Flush 'addr' from the kernel PCID: */
    invlpg(addr);

    /* If PTI is off there is no user PCID and nothing to flush. */
    if (!static_cpu_has(X86_FEATURE_PTI))
        return;

    loaded_mm_asid = this_cpu_read(cpu_tlbstate.loaded_mm_asid);
    cpu_pcide      = this_cpu_read(cpu_tlbstate.cr4) & X86_CR4_PCIDE;

    /*
     * invpcid_flush_one(pcid>0) will #GP if CR4.PCIDE==0. Check
     * 'cpu_pcide' to ensure that *this* CPU will not trigger those
     * #GP's even if called before CR4.PCIDE has been initialized.
     */
    if (boot_cpu_has(X86_FEATURE_INVPCID) && cpu_pcide)
        invpcid_flush_one(user_pcid(loaded_mm_asid), addr);
    else
        invalidate_user_asid(loaded_mm_asid);
}

per-cpu cpu_tlbstate.loaded_mm_asid      switch_mm_irqs_off mm_struct
this_cpu_write(cpu_tlbstate.loaded_mm_asid, ns.asid);

ns = choose_new_asid(next, next_tlb_gen);

this_cpu_write(cpu_tlbstate.loaded_mm_asid, ns.asid);

```

```

static struct new_asid choose_new_asid(struct mm_struct *next, u64 next_tlb_gen)
{
    struct new_asid ns;
    u16 asid;

    if (!static_cpu_has(X86_FEATURE_PCID)) {
        ns.asid = 0;
        ns.need_flush = 1;
        return ns;
    }

    /*
     * TLB consistency for global ASIDs is maintained with hardware assisted
     * remote TLB flushing. Global ASIDs are always up to date.
     */
    if (cpu_feature_enabled(X86_FEATURE_INVLPG)) {
        u16 global_asid = mm_global_asid(next);

        if (global_asid) {
            ns.asid = global_asid;
            ns.need_flush = 0;
            return ns;
        }
    }

    if (this_cpu_read(cpu_tlbstate.invalidate_other))
        clear_asid_other();

    for (asid = 0; asid < TLB_NR_DYN_ASIDS; asid++) {
        if (this_cpu_read(cpu_tlbstate.ctxs[asid].ctx_id) !=
            next->context.ctx_id)
            continue;

        ns.asid = asid;
        ns.need_flush = (this_cpu_read(cpu_tlbstate.ctxs[asid].tlb_gen) < next_tlb_gen);
        return ns;
    }

    /*
     * We don't currently own an ASID slot on this CPU.
     * Allocate a slot.
     */
    ns.asid = this_cpu_add_return(cpu_tlbstate.next_asid, 1) - 1;
    if (ns.asid >= TLB_NR_DYN_ASIDS) {
        ns.asid = 0;
        this_cpu_write(cpu_tlbstate.next_asid, 1);
    }
}

```

```

    }

    ns.need_flush = true;

    return ns;
}

asid pcid      switch_mm_irqs_off->load_new_mm_cr3->write_cr3      cr3      tlb_state per-
cpu      cpu tlb      tlb_state::loaded_mm_asid      tlb_state::ctxs type: tlb_context      tlb_context:::
mm_struct      mm mm tlb      copy_from_user()/copy_to_user()      pcid 0
>PA      PCID=0 TLB      B      PCID=0      TLB      A      B      TLB
PCID Linux      PCID user_pcid = base_pcid | (1 «11) ernel_pcid =
base_pcid base_pcid      5      cr3      pcid tlb      CONFIG_MITIGATION_PAGE_TABLE_ISOLATION
fault PTE      CPL=0      (CPL=0)      PTI      Meltdown      CPU
Intel (out-of-order execution)      CPU      speculatively      (cache)
pcid/asid      error_entry      IBRS_ENTER

/*
 * IBRS kernel mitigation for Spectre_v2.
 *
 * Assumes full context is established (PUSH_REGS, CR3 and GS) and it clobbers
 * the regs it uses (AX, CX, DX). Must be called before the first RET
 * instruction (NOTE! UNTRAIN_RET includes a RET instruction)
 *
 * The optional argument is used to save/restore the current value,
 * which is used on the paranoid paths.
 *
 * Assumes x86_spec_ctrl_{base,current} to have SPEC_CTRL_IBRS set.
*/
.macro IBRS_ENTER save_reg
#endif CONFIG_MITIGATION_IBRS_ENTRY
ALTERNATIVE "jmp .Lend_\@", "", X86_FEATURE_KERNEL_IBRS
movl $MSR_IA32_SPEC_CTRL, %ecx

.ifnb \save_reg
rdmsr
shl $32, %rdx
or %rdx, %rax
mov %rax, \save_reg
test $SPEC_CTRL_IBRS, %eax
jz .Ldo_wrmsr_\@
lfence
jmp .Lend_\@
.Ldo_wrmsr_\@:
.endif

```

```

    movq PER_CPU_VAR(x86_spec_ctrl_current), %rdx
    movl %edx, %eax
    shr $32, %rdx
    wrmsr
.Lend_\@:
#endif
.endm

MSR_IA32_SPEC_CTRL 0x48 intel sdm vol4 Register Address 0x48 msr Speculation
Control Bit 0 IBRS Indirect Branch Restricted Speculation 1 CPU
v2 / wrmsr CPU Enhanced IBRS IBRS MSR
wrmsr msr ecx edx:eax 32 64
UNTRAIN_RET_FROM_CALL
error_entry

leaq 8(%rsp), %rdi /* arg0 = pt_regs pointer */
/* Put us onto the real thread stack. */
jmp sync_regs

rsp 8 pt_regs rsp error_entry asm_common_interrupt call error_entry
mov ffffff81e0154d +8 r15

ffffffff81e01540 <asm_common_interrupt>:
fffffff81e01540: f3 0f 1e fa endbr64
fffffff81e01544: 90 nop
fffffff81e01545: 90 nop
fffffff81e01546: 90 nop
fffffff81e01547: fc cld
fffffff81e01548: e8 e3 06 00 00 call ffffff81e01c30 <error_entry>
fffffff81e0154d: 48 89 c4 mov %rax,%rsp
fffffff81e01550: 48 89 e7 mov %rsp,%rdi
fffffff81e01553: 48 8b 74 24 78 mov 0x78(%rsp),%rsi
fffffff81e01558: 48 c7 44 24 78 ff ff movq $0xfffffffffffffff,0x78(%rsp)
fffffff81e0155f: ff ff
fffffff81e01561: e8 5a 0c ef ff call ffffff81cf21c0 <common_interrupt>
fffffff81e01566: e9 05 08 00 00 jmp ffffff81e01d70 <error_return>
fffffff81e0156b: 66 66 2e 0f 1f 84 00 data16 cs nopw 0x0(%rax,%rax,1)
fffffff81e01572: 00 00 00 00
fffffff81e01576: 66 2e 0f 1f 84 00 00 cs nopw 0x0(%rax,%rax,1)
fffffff81e0157d: 00 00 00

jmp sync_regs sync_regs ret asm_common_interrupt sync_regs

/*
 * Help handler running on a per-cpu (IST or entry trampoline) stack
 * to switch to the normal thread stack if the interrupted code was in
 * user mode. The actual stack switch is done in entry_64.S

```

```

*/
asmlinkage __visible noinstr struct pt_regs *sync_regs(struct pt_regs *eregs)
{
    struct pt_regs *regs = (struct pt_regs *)current_top_of_stack() - 1;
    if (regs != eregs)
        *regs = *eregs;
    return regs;
}

                current_top_of_stack          task_struct::stack
error_entry           error_entry
sync_regs error_entry    rxax           pt_regs      /      asm_common_interrupt
idtentry_body

/**
 * idtentry_body - Macro to emit code calling the C function
 * @cfunc: C function to be called
 * @has_error_code: Hardware pushed error code on stack
 */
.macro idtentry_body cfunc has_error_code:req

/*
 * Call error_entry() and switch to the task stack if from userspace.
 *
 * When in XENPV, it is already in the task stack, and it can't fault
 * for native_iret() nor native_load_gs_index() since XENPV uses its
 * own pvops for IRET and load_gs_index(). And it doesn't need to
 * switch the CR3. So it can skip invoking error_entry().
 */
ALTERNATIVE "call error_entry; movq %rax, %rsp", \
            "call xen_error_entry", X86_FEATURE_XENPV

ENCODE_FRAME_POINTER
UNWIND_HINT_REGS

movq %rsp, %rdi /* pt_regs pointer into 1st argument*/

.if \has_error_code == 1
    movq ORIG_RAX(%rsp), %rsi /* get error code into 2nd argument*/
    movq $-1, ORIG_RAX(%rsp) /* no syscall to restart */
.endif

/* For some configurations \cfunc ends up being a noreturn. */
ANNOTATE_REACHABLE
call \cfunc

jmp error_return
.endm

```

```

error_entry->sync_regs ret idtentry_body rax      pt_regs    mov %rsp,
%rdi       rdi       rsi       pt_regs orig_rax      irq_entries_start push vector cfunc
>idtentry->idtentry_body common_interrupt call  c      " "
common_interrupt

/*
 * common_interrupt() handles all normal device IRQ's (the special SMP
 * cross-CPU interrupts have their own entry points).
 */
#define DEFINE_IDTENTRY_IRQ(func)
{
    struct pt_regs *old_regs = set_irq_regs(regs);

    /* entry code tells RCU that we're not quiescent. Check it. */
    RCU_LOCKDEP_WARN(!rcu_is_watching(), "IRQ failed to wake up RCU");

    if (unlikely(!call_irq_handler(vector, regs)))
        apic_eoi();

    set_irq_regs(old_regs);
}

common_interrupt     DEFINE_IDTENTRY_IRQ irq
#define DEFINE_IDTENTRY_IRQ(func)
    DEFINE_IDTENTRY_IRQ

/***
 * DEFINE_IDTENTRY_IRQ - Emit code for device interrupt IDT entry points
 * @func: Function name of the entry point
 *
 * The vector number is pushed by the low level entry stub and handed
 * to the function as error_code argument which needs to be truncated
 * to an u8 because the push is sign extending.
 *
 * irq_enter/exit_rcu() are invoked before the function body and the
 * KVM L1D flush request is set. Stack switching to the interrupt stack
 * has to be done in the function body if necessary.
 */
#define DEFINE_IDTENTRY_IRQ(func) \
static void __#func(struct pt_regs *regs, u32 vector); \
\ \
__visible noinstr void func(struct pt_regs *regs, \
    unsigned long error_code) \
{ \
    \
    irqentry_state_t state = irqentry_enter(regs); \
    u32 vector = (u32)(u8)error_code; \
}

```

```

        \
kvm_set_cpu_l1tf_flush_l1d();                                \
instrumentation_begin();          \
run_irq_on_irqstack_cond(_##func, regs, vector);  \
instrumentation_end();          \
irqentry_exit(regs, state);    \
}          \
        \
static noinline void _##func(struct pt_regs *regs, u32 vector)
common_interrupt

static void __common_interrupt(struct pt_regs *regs, u32 vector);

__visible noinstr void common_interrupt(struct pt_regs *regs,
unsigned long error_code)
{
    irqentry_state_t state = irqentry_enter(regs);
    u32 vector = (u32)(u8)error_code;

    kvm_set_cpu_l1tf_flush_l1d();
    instrumentation_begin();
    run_irq_on_irqstack_cond(__common_interrupt, regs, vector);
    instrumentation_end();
    irqentry_exit(regs, state);
}

static noinline void __common_interrupt(struct pt_regs *regs, u32 vector)
{
    struct pt_regs *old_regs = set_irq_regs(regs);

    /* entry code tells RCU that we're not quiescent. Check it. */
    RCU_LOCKDEP_WARN(!rcu_is_watching(), "IRQ failed to wake up RCU");

    if (unlikely(!call_irq_handler(vector, regs)))
        apic_eoi();

    set_irq_regs(old_regs);
}

irqentry_enter/irqentry_exit           irqentry      RCU lockdep
error_code      vector      u8   u32  vector
kvm_set_cpu_l1tf_flush_l1d      guest

arch/x86/include/asm/hardirq.h

static __always_inline void kvm_set_cpu_l1tf_flush_l1d(void)

```

```

{
    __this_cpu_write(irq_stat.kvm_cpu_l1tf_flush_l1d, 1);
}

guest vmx_l1d_flush          l1d cache
instrumentation_begin/instrumentation_end   CONFIG_NOINSTR_VALIDATION
run_irq_on_irqstack_cond

arch/x86/include/asm/irq_stack.h

#define run_irq_on_irqstack_cond(func, regs, vector) \
{ \
    assert_function_type(func, void (*)(struct pt_regs *, u32)); \
    assert_arg_type(regs, struct pt_regs *); \
    assert_arg_type(vector, u32); \
    \
    call_on_irqstack_cond(func, regs, ASM_CALL_IRQ, \
                          IRQ_CONSTRAINTS, regs, vector); \
}
}

arch/x86/include/asm/irq_stack.h

/*
 * Macro to invoke system vector and device interrupt C handlers.
 */
#define call_on_irqstack_cond(func, regs, asm_call, constr, c_args...) \
{ \
    /* \
     * User mode entry and interrupt on the irq stack do not \
     * switch stacks. If from user mode the task stack is empty. \
     */ \
    if (user_mode(regs) || __this_cpu_read(hardirq_stack_inuse)) { \
        irq_enter_rcu(); \
        func(c_args); \
        irq_exit_rcu(); \
    } else { \
        /* \
         * Mark the irq stack inuse _before_ and unmark _after_ \
         * switching stacks. Interrupts are disabled in both \
         * places. Invoke the stack switch macro with the call \
         * sequence which matches the above direct invocation. \
         */ \
        __this_cpu_write(hardirq_stack_inuse, true); \
        call_on_irqstack(func, asm_call, constr); \
        __this_cpu_write(hardirq_stack_inuse, false); \
    } \
}

```

```

common_interrupt           irq_entries_start/error_entry      irq_ent

IRQ

CPU
-      TSS.sp0  IST      ← "  "
-      SS, RSP, RFLAGS, CS, RIP
-      IDT

entry_64.S

    eregs
    sync_regs(eregs)

    eregs → task
    regs

    **  **

call_on_irqstack_cond(func, regs, ...)
    user_mode(regs) == true
        func()

func() =         handle_irq_event
    task

irq_exit_rcu() →

iret →
    rsp  task      CPU           call_on_irqstack_cond  else      task
    task

    CPL=0
    → CPU

(entry_64.S)
    task

```

```

call_on_irqstack_cond(func, regs, ...)
    user_mode(regs) == false
    hardirq_stack_inuse == false
        per-CPU hardirq
            call_on_irqstack(func, asm_call, ...)

__this_cpu_read(hardirq_stack_ptr)
    → map_irq_stack()    per-CPU

func() = handle_irq_event
    per-CPU **hardirq **

__this_cpu_write(hardirq_stack_inuse, false)
irq_exit_rcu()

iret
    hardirq_stack_inuse true          CPU  hardirq  hardirq_stack_inuse
                                    hardirq

→ CPU      CPL=0

entry_64.S

call_on_irqstack_cond(func, regs, ...)
    user_mode(regs) == false
    hardirq_stack_inuse == true

        func()    hardirq

func() = IRQ handler
    CPU  **  hardirq **

irq_exit_rcu() →
    “CPU      TSS.IST TSS.sp0      ”          sp0      tss task-
state segment 64      intel sdm vol3 Figure 7-11
    rsp0  rsp1  rsp2  7 IST Intel sdm vol3      call/interrupt/exception

```

When the processor performs a call to the exception- or interrupt-handler procedure:

- The segment selector and stack pointer for the stack to be used by the handler are obtained for the currently executing task. On this new stack, the processor pushes the stack segment stack pointer of the interrupted procedure.
- The processor then saves the current state of the EFLAGS, CS, and EIP registers on the ne

I/O Map Base Address	Reserved	0
Reserved	96	
Reserved	92	
IST7 (upper 32 bits)	88	
IST7 (lower 32 bits)	84	
IST6 (upper 32 bits)	80	
IST6 (lower 32 bits)	76	
IST5 (upper 32 bits)	72	
IST5 (lower 32 bits)	68	
IST4 (upper 32 bits)	64	
IST4 (lower 32 bits)	60	
IST3 (upper 32 bits)	56	
IST3 (lower 32 bits)	52	
IST2 (upper 32 bits)	48	
IST2 (lower 32 bits)	44	
IST1 (upper 32 bits)	40	
IST1 (lower 32 bits)	36	
Reserved	32	
Reserved	28	
RSP2 (upper 32 bits)	24	
RSP2 (lower 32 bits)	20	
RSP1 (upper 32 bits)	16	
RSP1 (lower 32 bits)	12	
RSP0 (upper 32 bits)	8	
RSP0 (lower 32 bits)	4	
Reserved	0	

[ ] Reserved bits. Set to 0.

Figure 3: 64bit tss

Figure 6-4).

c. If an exception causes an error code to be saved, it is pushed on the new stack after the

cpu	user mode 3	0	tss      tss    rsp   Intel
sdm vol3			

If the call requires a change in privilege level, the processor also switches to the stack f

rsp      level 0      rsp0	Linux	x86_hw_tss      x86-64
----------------------------	-------	------------------------

arch/x86/include/asm/processor.h

```

struct x86_hw_tss {
    u32 reserved1;
    u64 sp0;
    u64 sp1;

    /*
     * Since Linux does not use ring 2, the 'sp2' slot is unused by
     * hardware. entry_SYSCALL_64 uses it as scratch space to stash
     * the user RSP value.
     */
    u64 sp2;
}

```

```

u64 reserved2;
u64 ist[7];
u32 reserved3;
u32 reserved4;
u16 reserved5;
u16 io_bitmap_base;

} __attribute__((packed));

```

cpu	tss	tss	base	limit	x86-64	tss
-----	-----	-----	------	-------	--------	-----

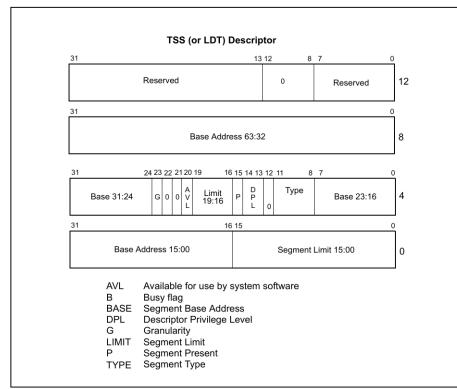


Figure 4: 64 tss/ldt

arch/x86/include/asm/desc\_defs.h

```

/* LDT or TSS descriptor in the GDT. */
struct ldttss_desc {
    u16 limit0;
    u16 base0;

    u16 base1 : 8, type : 5, dpl : 2, p : 1;
    u16 limit1 : 4, zero0 : 3, g : 1, base2 : 8;
#endif CONFIG_X86_64
    u32 base3;
    u32 zero1;
#endif
} __attribute__((packed));

```

x86-64	tss	sp0	sp1	sp2	tss
X86 cpu	16 real mode	32 protected mode	x86 cpu	64	IA-
32e mode	64-bit mode	Compatibility	64	32 app	

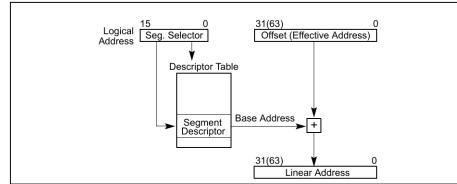


Figure 5:

```

intel sdm vol3
      GDT Global Descriptor Table LDT Local Descriptor Table
      GDTR Global Descriptor Table Register   GDT   LDTR Local Descrip-
      torTable Register   LDT           CS SS DS ES FS GS   GDT LDT

```

Visible Part	Hidden Part	
Segment Selector	Base Address, Limit, Access Information	CS
		SS
		DS
		ES
		FS
		GS

Figure 6:

gdt/ldt

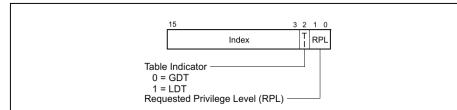


Figure 7:

```

bit 3-15      gdt/ldt      TI      gdt/ldt
      task register x86    tss    task register      tss    gdt      gdt tss
                                         TR
                                         gdt      lgdt      gdt tss      task
                                         gdt
                                         Linux   gdt      Linux   cpu      gdt page
                                         arch/x86/include/asm/cpu_entry_area.h

/*
 * cpu_entry_area is a percpu region that contains things needed by the CPU
 * and early entry/exit code. Real types aren't used for all fields here

```

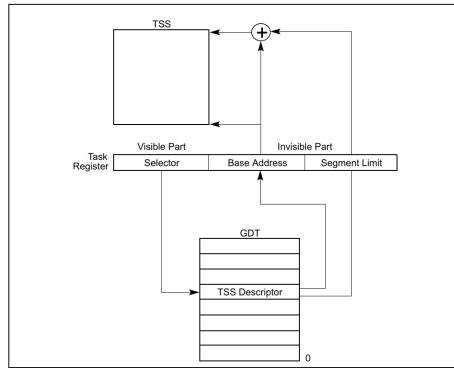


Figure 8: task register tss

```

* to avoid circular header dependencies.
*
* Every field is a virtual alias of some other allocated backing store.
* There is no direct allocation of a struct cpu_entry_area.
*/

```

```

struct cpu_entry_area {
    char gdt[PAGE_SIZE];
    ...
}

setup_cpu_entry_area      gdt page
arch/x86/mm/cpu_entry_area.c

/* Setup the fixmap mappings only once per-processor */
static void __init setup_cpu_entry_area(unsigned int cpu)
{
    struct cpu_entry_area *cea = get_cpu_entry_area(cpu);
#ifdef CONFIG_X86_64
    /* On 64-bit systems, we use a read-only fixmap GDT and TSS. */
    pgprot_t gdt_prot = PAGE_KERNEL_RO;
    pgprot_t tss_prot = PAGE_KERNEL_RO;
#else
    /*
     * On 32-bit systems, the GDT cannot be read-only because
     * our double fault handler uses a task gate, and entering through
     * a task gate needs to change an available TSS to busy. If the
     * GDT is read-only, that will triple fault. The TSS cannot be
     * read-only because the CPU writes to it on task switches.

```

```

        */
pgprot_t gdt_prot = PAGE_KERNEL;
pgprot_t tss_prot = PAGE_KERNEL;
#endif
        ...
cea_set_pte(&cea->gdt, get_cpu_gdt_paddr(cpu), gdt_prot);
        ...
}
cea_set_pte      cpu_entry_area           alloc_*      percpu      cea_set_pte      cpu_entry
arch/x86/include/asm/desc.h

/* Provide the physical address of the GDT page. */
static inline phys_addr_t get_cpu_gdt_paddr(unsigned int cpu)
{
    return per_cpu_ptr_to_phys(get_cpu_gdt_rw(cpu));
}

static inline struct desc_struct *get_cpu_gdt_rw(unsigned int cpu)
{
    return per_cpu(gdt_page, cpu).gdt;
}

cpu_entry_area      backingstore      gdt_page per_cpu      percpu      DEFINE
arch/x86/kernel/cpu/common.c

DEFINE_PER_CPU_PAGE_ALIGNED(struct gdt_page, gdt_page) = { .gdt = {
#endif CONFIG_X86_64
/*
 * We need valid kernel segments for data and code in long mode too
 * IRET will check the segment types kkeil 2000/10/28
 * Also sysret mandates a special GDT layout
 *
 * TLS descriptors are currently at a different place compared to i386.
 * Hopefully nobody expects them at a fixed place (Wine?)
*/
[GDT_ENTRY_KERNEL32_CS] = GDT_ENTRY_INIT(DESC_CODE32, 0, 0xffff),
[GDT_ENTRY_KERNEL_CS] = GDT_ENTRY_INIT(DESC_CODE64, 0, 0xffff),
[GDT_ENTRY_KERNEL_DS] = GDT_ENTRY_INIT(DESC_DATA64, 0, 0xffff),
[GDT_ENTRY_DEFAULT_USER32_CS] = GDT_ENTRY_INIT(DESC_CODE32 | DESC_USER, 0, 0xffff),
[GDT_ENTRY_DEFAULT_USER_DS] = GDT_ENTRY_INIT(DESC_DATA64 | DESC_USER, 0, 0xffff),
[GDT_ENTRY_DEFAULT_USER_CS] = GDT_ENTRY_INIT(DESC_CODE64 | DESC_USER, 0, 0xffff),
#else
[GDT_ENTRY_KERNEL_CS] = GDT_ENTRY_INIT(DESC_CODE32, 0, 0xffff),
[GDT_ENTRY_KERNEL_DS] = GDT_ENTRY_INIT(DESC_DATA32, 0, 0xffff),

```

```

[GDT_ENTRY_DEFAULT_USER_CS] = GDT_ENTRY_INIT(DESC_CODE32 | DESC_USER, 0, 0xffff),
[GDT_ENTRY_DEFAULT_USER_DS] = GDT_ENTRY_INIT(DESC_DATA32 | DESC_USER, 0, 0xffff),
*/
 * Segments used for calling PnP BIOS have byte granularity.
 * They code segments and data segments have fixed 64k limits,
 * the transfer segment sizes are set at run time.
 */
[GDT_ENTRY_PNPBIOS_CS32] = GDT_ENTRY_INIT(DESC_CODE32_BIOS, 0, 0xffff),
[GDT_ENTRY_PNPBIOS_CS16] = GDT_ENTRY_INIT(DESC_CODE16, 0, 0xffff),
[GDT_ENTRY_PNPBIOS_DS] = GDT_ENTRY_INIT(DESC_DATA16, 0, 0xffff),
[GDT_ENTRY_PNPBIOS_TS1] = GDT_ENTRY_INIT(DESC_DATA16, 0, 0),
[GDT_ENTRY_PNPBIOS_TS2] = GDT_ENTRY_INIT(DESC_DATA16, 0, 0),
/*
 * The APM segments have byte granularity and their bases
 * are set at run time. All have 64k limits.
 */
[GDT_ENTRY_APMBIOS_BASE] = GDT_ENTRY_INIT(DESC_CODE32_BIOS, 0, 0xffff),
[GDT_ENTRY_APMBIOS_BASE+1] = GDT_ENTRY_INIT(DESC_CODE16, 0, 0xffff),
[GDT_ENTRY_APMBIOS_BASE+2] = GDT_ENTRY_INIT(DESC_DATA32_BIOS, 0, 0xffff),

[GDT_ENTRY_ESPFIX_SS] = GDT_ENTRY_INIT(DESC_DATA32, 0, 0xffff),
[GDT_ENTRY_PERCPU] = GDT_ENTRY_INIT(DESC_DATA32, 0, 0xffff),
#endif
}};

include/linux/percpu-defs.h

#define DEFINE_PER_CPU_PAGE_ALIGNED(type, name) \
DEFINE_PER_CPU_SECTION(type, name, "..page_aligned") \
__aligned(PAGE_SIZE)

#define DEFINE_PER_CPU_SECTION(type, name, sec) \
__PCPU_ATTRS(sec) __typeof__(type) name

/*
 * Base implementations of per-CPU variable declarations and definitions, where
 * the section in which the variable is to be placed is provided by the
 * 'sec' argument. This may be used to affect the parameters governing the
 * variable's storage.
 *
 * NOTE! The sections for the DECLARE and for the DEFINE must match, lest
 * linkage errors occur due to the compiler generating the wrong code to access
 * that section.
 */
#define __PCPU_ATTRS(sec) \

```

```

__percpu __attribute__((section(PER_CPU_BASE_SECTION sec))) \
PER_CPU_ATTRIBUTES

#ifndef PER_CPU_BASE_SECTION
#define CONFIG_SMP
#define PER_CPU_BASE_SECTION ".data..percpu"

    percpu          section .data..percpu.page_aligned section percpu      vmlinux
-S .data.percpu..page_aligned section .data..percpu section      percpu section .data..percpu
section

include/asm-generic/vmlinux.lds.h

/***
 * PERCPU_INPUT - the percpu input sections
 * @cacheline: cacheline size
 *
 * The core percpu section names and core symbols which do not rely
 * directly upon load addresses.
 *
 * @cacheline is used to align subsections to avoid false cacheline
 * sharing between subsections for different purposes.
 */
#define PERCPU_INPUT(cacheline)      \
__per_cpu_start = .;      \
. = ALIGN(PAGE_SIZE);      \
*(.data..percpu..page_aligned)      \
. = ALIGN(cacheline);      \
__per_cpu_hot_start = .;      \
*(SORT_BY_ALIGNMENT(.data..percpu..hot.*))      \
__per_cpu_hot_end = .;      \
. = ALIGN(cacheline);      \
*(.data..percpu..read_mostly)      \
. = ALIGN(cacheline);      \
*(.data..percpu)      \
*(.data..percpu..shared_aligned)      \
PERCPU_DECRYPTED_SECTION      \
__per_cpu_end = .;

/***
 * PERCPU_SECTION - define output section for percpu area
 * @cacheline: cacheline size
 *
 * Macro which expands to output section for percpu area.
 *
 * @cacheline is used to align subsections to avoid false cacheline
 * sharing between subsections for different purposes.
*/

```

```

*/
#define PERCPU_SECTION(cacheline)      \
    . = ALIGN(PAGE_SIZE);           \
    .data..percpu : AT(ADDR(.data..percpu) - LOAD_OFFSET) { \
        PERCPU_INPUT(cacheline)      \
    }

    percpu section .data..percpu PERCPU_SECTION->PERCPU_INPUT section __per_c
section

arch/x86/kernel/vmlinux.lds.S

PERCPU_SECTION(L1_CACHE_BYTES)

.data..percpu section Type PROGBITS Flags WA vmlinux bootloader
offset      __per_cpu_start   percpu   percpu   percpu
            cea           setup_cpu_entry_area->get_cpu_entry_area

arch/x86/mm/cpu_entry_area.c

/* Is called from entry code, so must be noinstr */
noinstr struct cpu_entry_area *get_cpu_entry_area(int cpu)
{
    unsigned long va = CPU_ENTRY_AREA_PER_CPU + cea_offset(cpu) * CPU_ENTRY_AREA_SIZE;
    BUILD_BUG_ON(sizeof(struct cpu_entry_area) % PAGE_SIZE != 0);

    return (struct cpu_entry_area *) va;
}

arch/x86/include/asm/pgtable_areas.h

/* Single page reserved for the readonly IDT mapping: */
#define CPU_ENTRY_AREA_RO_IDT CPU_ENTRY_AREA_BASE
#define CPU_ENTRY_AREA_PER_CPU (CPU_ENTRY_AREA_RO_IDT + PAGE_SIZE)

CPU_ENTRY_AREA_BASE          CPU_ENTRY_AREA_BASE      entry
area per cpu     CPU_ENTRY_AREA_PER_CPU   cea_set_pte   per
cpu         percpu      cpu       cea_offset
            cea_set_pte   per cpu gdt_page cea->gdt      PAGE_KERNEL_RO      x86-
64     percpu gdt page      percpu back store page direct rw
            cea->gdt   fixmap ro   gdt page

arch/x86/include/asm/desc.h

/* Provide the original GDT */
static inline struct desc_struct *get_cpu_gdt_rw(unsigned int cpu)
{
    return per_cpu(gdt_page, cpu).gdt;
}

```

```

gdt page
arch/x86/include/asm/desc.h

/* Provide the fixmap address of the remapped GDT */
static inline struct desc_struct *get_cpu_gdt_rw(int cpu)
{
    return (struct desc_struct *)&get_cpu_entry_area(cpu)->gdt;
}

gdt page      gdtr
arch/x86/kernel/cpu/common.c

/* Load the original GDT from the per-cpu structure */
void load_direct_gdt(int cpu)
{
    struct desc_ptr gdt_descr;

    gdt_descr.address = (long)get_cpu_gdt_rw(cpu);
    gdt_descr.size = GDT_SIZE - 1;
    load_gdt(&gdt_descr);
}
EXPORT_SYMBOL_FOR_KVM(load_direct_gdt);

/* Load a fixmap remapping of the per-cpu GDT */
void load_fixmap_gdt(int cpu)
{
    struct desc_ptr gdt_descr;

    gdt_descr.address = (long)get_cpu_gdt_ro(cpu);
    gdt_descr.size = GDT_SIZE - 1;
    load_gdt(&gdt_descr);
}
EXPORT_SYMBOL_GPL(load_fixmap_gdt);

gdtr      load_gdt  lgdt   gdtr
arch/x86/include/asm/desc.h

#define load_gdt(dtr)    native_load_gdt(dtr)

static inline void native_load_gdt(const struct desc_ptr *dtr)
{
    asm volatile("lgdt %0::"m" (*dtr));
}

Linux    gdt

```

```

Linux tss      Linux  cpu_init_exception_handling    set_tss_desc tss      load_TR_desc tss
register TR

arch/x86/kernel/cpu/common.c

/*
 * Setup everything needed to handle exceptions from the IDT, including the IST
 * exceptions which use paranoid_entry().
 */
void cpu_init_exception_handling(bool boot_cpu)
{
    ...
    set_tss_desc(cpu, &get_cpu_entry_area(cpu)->tss.x86_tss);

    load_TR_desc();
    ...
}

set_tss_desc          tss.x86_tss      gdt_page      cpu_entry_area      percpu      backing
store   setup_cpu_entry_area

/* Setup the fixmap mappings only once per-processor */
static void __init setup_cpu_entry_area(unsigned int cpu)
{
    struct cpu_entry_area *cea = get_cpu_entry_area(cpu);
    pgprot_t tss_prot = PAGE_KERNEL_RO;

    ...
    cea_map_percpu_pages(&cea->tss, &per_cpu(cpu_tss_rw, cpu),
                         sizeof(struct tss_struct) / PAGE_SIZE, tss_prot);
    ...
}

gdt tss  Linux      tss           rsp0

```